


Graphical Abstract

TEMPO: Regime-Aware Operator Learning with Locally Adapted POD Bases

Anonymous Authors



Highlights

TEMPO: Regime-Aware Operator Learning with Locally Adapted POD Bases

Anonymous Authors

- TEMPO discovers solution regimes from data and fits a dedicated basis to each.
- Regime-local bases match or outperform a single global basis of the same rank.
- TEMPO(POD) reduces error by 13–50% on Burgers’ and up to $4.75\times$ on Darcy flow.
- TEMPO trains across all parameter values and handles multiple solution regimes.
- Regime structure reflects initial conditions and physical parameters without labels.

TEMPO: Regime-Aware Operator Learning with Locally Adapted POD Bases

Anonymous Authors

Anonymous Institution,

Abstract

Operator learning approximates the solution map of a parametric PDE from data, enabling fast prediction at new parameter values without repeated simulation. The dominant approach fits a single global low-rank basis to all solution snapshots — an assumption that fails when the governing parameter drives qualitatively distinct solution regimes, such as viscosity separating smooth diffusive profiles from near-discontinuous shocks.

We propose TEMPO (Trajectory EM Mixture of POD Operators), a two-phase framework for regime-aware operator learning. In the offline phase, TEMPO recovers latent regimes via Expectation-Maximization, fitting a dedicated linear or nonlinear basis to each discovered regime. We prove that regime-local bases achieve reconstruction error no greater than a single global basis of the same rank. In the online phase, a gating network and per-regime branch networks are trained jointly, routing new inputs to the correct regime at inference.

TEMPO(POD) reduces error by 13–50% over the POD-DeepONet baseline on 1D Burgers’ equation. On 2D Darcy flow, per-parameter specialists fail out-of-distribution, while TEMPO(POD) reduces error by up to $4.75\times$. The discovered regime structure reflects both initial-condition shape and the physical parameter, recovered from trajectory data without regime-label supervision.

Keywords: operator learning, parametric PDE, regime discovery, proper orthogonal decomposition, expectation-maximization, neural operator

1. Introduction

Reduced-order modeling of parametric PDEs constructs low-rank bases from solution snapshots to enable fast prediction at new parameter values without resolving the governing equations [1, 2]. The dominant approach, exemplified by POD-DeepONet [3], fits a single global basis to the full snapshot ensemble and trains a branch network to predict expansion coefficients for unseen inputs [4, 5]. This approach succeeds when solutions lie on a smooth low-dimensional manifold, but fails when the governing parameter drives qualitatively distinct dynamics: viscosity separating laminar flow from near-discontinuous shocks [6],

or source magnitude shifting solution amplitudes by orders of magnitude. A single global basis must distribute its capacity across all regimes, representing none faithfully.

Partitioning the parameter space and fitting local bases to each region is the classical remedy in reduced-order modeling [7, 8]. Daniel et al. [9] extend this to deep learning by clustering solution subspaces on the Grassmann manifold and training a dictionary selector at inference. These approaches reduce approximation error on multi-modal problems, but share a fundamental limitation: clustering is performed independently of basis quality, so the discovered partition need not align with the directions of greatest reconstruction benefit.

A second line of work enriches the basis while retaining a single global model. POD-PoU [10] blends several POD bases via a partition-of-unity function, reducing error on sharp-gradient problems. Nonlinear ROM using deep autoencoders [11] has demonstrated that low-dimensional nonlinear manifolds can dramatically outperform linear subspaces. NeuralPOD [12] pursues this direction for operator learning, learning nonlinear mode functions by greedy residual minimisation. These approaches improve representational capacity but treat the full ensemble jointly, without identifying distinct solution regimes.

Neural operators such as FNO [13] and CNO [14] learn solution operators as discretisation-invariant functional maps, achieving high accuracy without explicit basis decompositions. Their purpose is complementary to the present work: they do not produce inspectable basis representations or identify latent regime structure.

We propose **TEMPO** (*Trajectory EM Mixture of POD Operators*), which addresses this by coupling regime discovery directly with basis quality in a principled probabilistic framework. TEMPO models the trajectory ensemble as a Gaussian mixture and recovers regime structure via Expectation-Maximization (EM) [15]. The E-step assigns each trajectory to regimes based on reconstruction quality; the M-step fits a dedicated basis to each regime. Basis quality and regime assignments improve each other iteratively. After convergence, a GatingNet and per-regime BranchNets are trained jointly, routing new inputs to the correct regime at inference.

Contributions:

- **TEMPO framework.** A two-phase mixture-of-operators for parametric PDEs: an offline EM that discovers M latent regimes from unlabelled snapshots and fits a dedicated orthogonal (POD) or nonlinear (NeuralPOD) basis to each; an online phase where a GatingNet and M BranchNets are trained with a KL penalty that anchors routing to the offline regime partition.
- **Adaptive basis update.** A distribution-shift criterion that monitors moment changes in regime composition across EM iterations and selects per-regime between full retraining, incremental mode addition/pruning, or skipping.
- **Theoretical guarantee.** Regime-local bases achieve reconstruction error no greater than the global rank- K optimum for any soft-assignment matrix

54 and any M, K (Proposition 1).

- 55 • **Experiments.** On 1D Burgers', TEMPO(POD) reduces error by 13–50%
 56 over the jointly-trained POD-DeepONet baseline; on 2D Darcy flow, error
 57 is reduced by up to $4.75\times$, while per-parameter specialists fail on unseen
 58 parameter values.

59 2. Problem Statement

60 Let $\Omega \subset \mathbb{R}^d$ be a bounded spatial domain and $\mathcal{T} \subset \mathbb{R}_{\geq 0}$ a time interval.
 61 We are given a dataset of N solution trajectories:

$$\mathcal{D} = \{ (u_0^{(i)}, \kappa^{(i)}, s^{(i)}) \}_{i=1}^N, \quad s^{(i)} = s(\cdot, \cdot; u_0^{(i)}, \kappa^{(i)}) \in L^2(\Omega \times \mathcal{T}), \quad (1)$$

62 where $u_0^{(i)} \in \mathcal{U} = L^2(\Omega)$ is the initial condition and $\kappa^{(i)} \in \mathcal{K}$ the physical
 63 parameter. Each trajectory $s^{(i)}$ is discretized on $N_y = N_t N_x$ quadrature points
 64 $\{y_j = (x_j, t_j), w_j\}_{j=1}^{N_y}$ covering the product grid $\Omega \times \mathcal{T}$, with weighted norm
 65 $\|f\|_w^2 = \sum_j w_j f(y_j)^2$.




Figure 1: Training dataset $\mathcal{D}$: solution trajectories $s^{(i)}$ on $\Omega \times \mathcal{T}$ for different $\kappa^{(i)} \in \mathcal{K}$.

66 **Definition 1** (Solution operator). *The solution operator*

$$\mathcal{G} : \mathcal{U} \times \mathcal{K} \rightarrow L^2(\Omega \times \mathcal{T})$$

67 *maps each pair (u_0, κ) to the full space-time solution: $\mathcal{G}(u_0, \kappa) = s(\cdot, \cdot; u_0, \kappa)$.*

68 The operator $\mathcal{G}$ is unknown and expensive to evaluate directly. The *operator*
 69 *learning problem* is to construct $\hat{\mathcal{G}} \approx \mathcal{G}$ from $\mathcal{D}$ alone, generalizing to unseen
 70 $(u_0, \kappa) \notin \mathcal{D}$ without solving the PDE.

71 *Multi-regime assumption.* When solutions change character across $\mathcal{K}$, the opti-
 72 mal low-dimensional representation differs from regime to regime, and no single
 73 global basis captures the full solution diversity. We formalize this observation
 74 via a latent variable model over the trajectory ensemble.

75 **Definition 2** (Generative model). *Each trajectory arises from one of M latent*
 76 *regimes. The regime label $z_i \in \{1, \dots, M\}$ is unobserved and drawn from a*
 77 *categorical prior,*

$$z_i \sim \text{Categorical}(\pi_1, \dots, \pi_M), \quad \pi_m \geq 0, \quad \sum_m \pi_m = 1. \quad (2)$$

78 *Conditioned on $z_i = m$, the trajectory is distributed as*

$$(s^{(i)} \mid z_i = m) \sim \mathcal{N}(\hat{s}_m^{(i)}, \sigma^2 W^{-1}), \quad (3)$$

79 *where $W = \text{diag}(w_1, \dots, w_{N_y})$ is the quadrature weight matrix, $\sigma^2 > 0$ controls*
 80 *the softness of regime assignments, and $\hat{s}_m^{(i)}$ is the regime- m approximation to*
 81 *$s^{(i)}$.*

82 The generative model leaves $\hat{s}_m^{(i)}$ unspecified. We require only that it admits
 83 a low-rank trajectory decomposition:

$$\hat{s}_m^{(i)}(y) = \phi_0^{(m)}(y) + \sum_{k=1}^{K_m} \lambda_k^{(m,i)} \phi_k^{(m)}(y), \quad y = (x, t) \in \Omega \times \mathcal{T}, \quad (4)$$

84 where $\phi_k^{(m)} : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$ are spatiotemporal mode functions and $\lambda_k^{(m,i)} \in \mathbb{R}$
 85 a scalar amplitude for trajectory i on mode k . The linear case recovers POD-
 86 DeepONet [3]; the nonlinear case recovers NeuralPOD [12]. Both are studied in
 87 this work.

88 **Definition 3** (Regime approximator). *The parameters of regime m are $\Phi_m =$*
 89 *$\{\phi_k^{(m)}\}_{k=0}^{K_m}$ and scalar amplitudes $\Lambda_m = \{\lambda_k^{(m,i)}\}_{k=1, i=1}^{K_m, N}$, where $\phi_0^{(m)} : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$*
 90 *is the regime mean field, $\{\phi_k^{(m)}\}_{k \geq 1}$ spatiotemporal mode functions, $\lambda_k^{(m,i)} \in \mathbb{R}$*
 91 *the scalar amplitude of trajectory i on mode k , and $K_m \leq K_{\max}$ the adaptive*
 92 *mode count.*

93 The learning problem decomposes into two phases: discovering the latent
 94 regime structure from $\mathcal{D}$ *offline*, then training a deployable operator for fast
 95 inference *online*.

96 *Phase 1: Offline regime discovery.* Since the regime labels z_i in (3) are unob-
 97 served, we recover the bases $\{\Phi_m\}$, mixture weights π , and noise level σ^2 by
 98 maximizing the observed-data log-likelihood:

$$(\Phi^*, \pi^*, \sigma^{2*}) = \arg \max_{\Phi, \pi, \sigma^2} \sum_{i=1}^N \log \sum_{m=1}^M \frac{\pi_m \det(W)^{1/2}}{(2\pi\sigma^2)^{N_y/2}} \exp \left(- \frac{\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2}{2\sigma^2} \right). \quad (5)$$

99 The solution yields soft regime responsibilities $\gamma_{im}^* = p(z_i = m \mid s^{(i)}, \Phi^*, \pi^*, \sigma^{2*})$
 100 that serve as supervision for Phase 2.

101 *Phase 2: Online operator learning.* With bases $\{\Phi_m^*\}$ and responsibilities $\{\gamma_{im}^*\}$
 102 frozen, we train a gating network (*GatingNet*) and M branch networks (*Branch-*
 103 *Nets*) that map a new input (u_0, κ) to a weighted combination of local regime
 104 predictions. The precise architecture and loss are given in Section 3.4. The
 105 quality of the final operator is measured by the mean relative L^2 error:

$$\varepsilon = \frac{1}{N_{\text{test}}} \sum_{i=1}^{N_{\text{test}}} \frac{\|s^{(i)} - \hat{s}(\cdot; u_0^{(i)}, \kappa^{(i)})\|_w}{\|s^{(i)}\|_w}. \quad (6)$$

106 3. Method

107 All methods studied in this work share a common two-phase structure: an
 108 *offline* basis-fitting phase that constructs a low-rank approximation to the so-
 109 lution manifold from snapshots in $\mathcal{D}$, and an *online* operator-learning phase
 110 that trains a branch network to predict expansion coefficients from a new input
 111 (u_0, κ) . The four methods differ along two independent axes:

- 112 • **Basis type.** A *linear* (POD) basis uses SVD of the (weighted) snap-
 113 shot matrix; a *nonlinear* (NeuralPOD) basis learns spatiotemporal mode
 114 networks by sequential residual minimization.
- 115 • **Number of regimes.** A *single-regime* operator fits one global basis to all
 116 snapshots; a *multi-regime* operator (TEMPO) discovers M latent regimes
 117 via EM and assigns a dedicated basis to each.

118 This gives four combinations: POD-DeepONet [3] (linear, global), NeuralPOD-
 119 DeepONet (nonlinear, global, our extension of [12]), TEMPO(POD) (linear, M
 120 regimes), and TEMPO(NeuralPOD) (nonlinear, M regimes).

121 3.1. Basis Representations

122 3.1.1. POD Basis

123 Proper Orthogonal Decomposition (POD) [16] extracts a hierarchy of or-
 124 thogonal modes that capture the maximum energy from the snapshot ensemble.
 125 Let $\phi_0 := \bar{s} = \frac{1}{N} \sum_i s^{(i)}$ be the mean field and $\tilde{S} \in \mathbb{R}^{N_y \times N}$ the matrix with
 126 columns $\tilde{s}^{(i)} = s^{(i)} - \phi_0$. The rank- K truncated SVD $\tilde{S} \approx U_K \Sigma_K V_K^\top$ yields
 127 orthonormal modes $\{\phi_k\}_{k=1}^K$ (columns of U_K) and coefficients $\lambda_k^{(i)} = [\Sigma_K V_K^\top]_{ki}$;
 128 the approximation

$$\hat{s}^{(i)} = \phi_0 + \sum_{k=1}^K \lambda_k^{(i)} \phi_k \quad (7)$$

129 is optimal in the L^2 sense among all rank- K affine representations.

130 For TEMPO(POD), the M-step (17) is solved analytically: set $\phi_0^{(m)} := \bar{s}_m =$
 131 $\frac{1}{N_m} \sum_i \gamma_{im} s^{(i)}$, where $N_m = \sum_i \gamma_{im}$, and apply SVD to the weighted centered
 132 matrix with columns $\sqrt{\gamma_{im}} (s^{(i)} - \phi_0^{(m)})$. This is the closed-form linear analogue
 133 of Algorithm 1.

3.1.2. NeuralPOD Basis

To overcome the linearity constraint, we parametrize each mode function $\phi_k : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$ as a Fourier feature network [17]; we call the resulting basis a *Fourier NeuralPOD* basis:

$$\phi_k(y) = f_{\theta_k}(\psi(y)), \quad \psi(y) = [\sin(2\pi By), \cos(2\pi By)], \quad (8)$$

where $B \in \mathbb{R}^{F \times d}$ is a fixed random matrix drawn once at initialization and $f_{\theta_k} : \mathbb{R}^{2F} \rightarrow \mathbb{R}$ is a small MLP with Tanh activations. Plain MLPs exhibit spectral bias [18]; the Fourier encoding resolves this. For multi-scale coverage, the rows of B are partitioned evenly across L frequency scales $\sigma_1 < \dots < \sigma_L$: the ℓ -th block is drawn from $\mathcal{N}(0, \sigma_\ell^2 I)$. The mean field ϕ_0 uses the same architecture.

The scalar amplitudes $\lambda_k^{(i)} \in \mathbb{R}$ are direct learnable parameters, one per training trajectory, optimized jointly with ϕ_k . After Phase 1 they are frozen and serve as regression targets for the BranchNet in Phase 2.

Modes are learned by *greedy residual minimization* [12]. Let $r^{(0,i)} = s^{(i)} - \phi_0$, where ϕ_0 is fitted to the snapshot mean $\bar{s}$ defined above. For each subsequent mode $k \geq 1$, given the residual

$$r^{(k-1,i)}(y) = s^{(i)}(y) - \phi_0(y) - \sum_{\ell=1}^{k-1} \lambda_\ell^{(i)} \phi_\ell(y), \quad (9)$$

we solve

$$\min_{\phi_k, \{\lambda_k^{(i)}\}} \sum_{i=1}^N \sum_j w_j \left(\lambda_k^{(i)} \phi_k(y_j) - r_j^{(k-1,i)} \right)^2 \quad (10)$$

(the global case with uniform weights; the γ -weighted generalization used in TEMPO is Algorithm 1). Mode addition stops when the residual falls below fraction τ of the initial variance,

$$\sum_{i=1}^N \|r^{(k,i)}\|_w^2 < \tau \cdot \sum_{i=1}^N \|s^{(i)} - \bar{s}\|_w^2, \quad (11)$$

or $k = K_{\max}$. Each mode captures the largest remaining variance component — the nonlinear analogue of Gram-Schmidt orthogonalization.

3.2. Single-Regime Operators

3.2.1. POD-DeepONet

POD-DeepONet [3] replaces the learned trunk network of DeepONet [4] with the fixed POD basis. After computing the global modes $\{\phi_k\}_{k=1}^K$ and training coefficients $\{\lambda^{(i)}\}$ offline, a branch network $\mathcal{B}_\theta : \mathbb{R}^{m_s+d_\kappa} \rightarrow \mathbb{R}^K$ is trained online by regression against these coefficients:

$$\mathcal{L}_{\text{branch}} = \frac{1}{N} \sum_{i=1}^N \|\mathcal{B}_\theta(u_0^{(i)}, \kappa^{(i)}) - \lambda^{(i)}\|^2. \quad (12)$$

162 The prediction at any $(x, t) \in \Omega \times \mathcal{T}$ is

$$\hat{s}(x, t; u_0, \kappa) = \phi_0(x, t) + \sum_{k=1}^K \mathcal{B}_{\theta, k}(u_0, \kappa) \phi_k(x, t). \quad (13)$$

163 3.2.2. *NeuralPOD-DeepONet*

164 Mou et al. [12] introduce NeuralPOD and pair it with a DeepONet branch
165 network to predict mode coefficients for unseen inputs. We adopt this combi-
166 nation, which we refer to as *NeuralPOD-DeepONet*, as a single-regime baseline
167 and as the per-regime basis learner within TEMPO(NeuralPOD).

168 The procedure follows POD-DeepONet with the Fourier NeuralPOD basis
169 in place of POD. Offline fitting yields coefficients $\lambda^{(i)} \in \mathbb{R}^K$; online, a branch
170 network minimizes (12) and prediction follows (13). At equal mode count K ,
171 the nonlinear basis spans a strictly richer function space than POD.

172 3.3. *TEMPO: Offline Regime Discovery*

173 A single global basis, whether linear or nonlinear, must compromise between
174 qualitatively distinct regimes, as the optimal basis geometry differs fundamen-
175 tally between them. TEMPO avoids this by constructing a dedicated basis for
176 each of M latent regimes, discovered from unlabelled trajectories via EM on the
177 generative model of Definition 2.

178 *Initialization.* A global POD is computed on all N trajectories; each trajec-
179 tory is projected to a d_0 -dimensional coefficient vector $\alpha^{(i)} \in \mathbb{R}^{d_0}$. Running a
180 standard GMM on $\{\alpha^{(i)}\}$ yields initial responsibilities $\gamma_{im}^{(0)}$, weights $\pi_m^{(0)}$, and
181 per-regime statistics $(\mu_m^{(0)}, \Sigma_m^{(0)})$. Each regime basis $\Phi_m^{(0)}$ is then fitted to the
182 responsibility-weighted ensemble: for TEMPO(NeuralPOD) via NEURALBASIS
183 (Algorithm 1); for TEMPO(POD) via the weighted SVD of Section 3.1. The
184 noise level σ^2 is calibrated from the initial reconstructions rather than treated
185 as a free hyperparameter:

$$\sigma^2 \leftarrow \frac{1}{2N} \sum_{i=1}^N \sum_{m=1}^M \gamma_{im}^{(0)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2, \quad (14)$$

186 σ^2 is held fixed throughout EM; improving bases reduce residuals, which pro-
187 gressively sharpen regime assignments.

188 *E-step.* At iteration q , the responsibility of trajectory i for regime m is updated
189 by Bayes' rule on the Gaussian likelihood (3):

$$\gamma_{im}^{(q)} = \frac{\pi_m^{(q-1)} \exp\left(-\frac{\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2}{2\sigma^2}\right)}{\sum_{j=1}^M \pi_j^{(q-1)} \exp\left(-\frac{\|s^{(i)} - \hat{s}_j^{(i)}\|_w^2}{2\sigma^2}\right)}, \quad (15)$$

190 where $\hat{s}_m^{(i)}$ is evaluated via (4) using $\Phi_m^{(q-1)}$ and $\Lambda_m^{(q-1)}$. As the bases improve
 191 in the M-step, the reconstructions $\hat{s}_m^{(i)}$ sharpen and the E-step automatically
 192 refines regime assignments in response — distinguishing TEMPO from a GMM
 193 applied to a fixed feature space. When κ drives large amplitude variation (as
 194 with $\beta \in [0.01, 100]$ in the Darcy experiments), replacing the squared norm
 195 with the relative squared norm $\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2 / \|s^{(i)}\|_w^2$ throughout (in (15), (14),
 196 and (17)) keeps E- and M-steps consistent under a heteroscedastic likelihood
 197 with noise variance proportional to $\|s^{(i)}\|_w^2$. When κ is observed, the GMM
 198 initialization can operate in $\log \kappa$ space, and an optional κ -prior term in the
 199 E-step anchors assignments to the known parameter structure.

200 *M-step.* With responsibilities fixed, the M-step gives a closed-form update for
 201 the mixture weights,

$$\pi_m^{(q)} \leftarrow \frac{N_m^{(q)}}{N}, \quad N_m^{(q)} = \sum_{i=1}^N \gamma_{im}^{(q)}, \quad (16)$$

202 and a weighted reconstruction problem for each basis:

$$\min_{\Phi_m, \Lambda_m} \sum_{i=1}^N \gamma_{im}^{(q)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2, \quad (17)$$

203 solved by weighted SVD for TEMPO(POD) or NEURALBASIS (Algorithm 1)
 204 for TEMPO(NeuralPOD). Since gradient descent solves the NeuralPOD sub-
 205 problem without guaranteeing global optimality, TEMPO is a Generalized EM
 206 algorithm [19], which guarantees that the observed-data log-likelihood is non-
 207 decreasing at every iteration.

208 **Proposition 1** (Regime Approximation Advantage). *For any $\{\gamma_{im}\} \geq 0$ with*
 209 *$\sum_m \gamma_{im} = 1$, any $M \geq 1$, any $K \geq 1$:*

$$\varepsilon_{\text{TEMPO}}^K \leq \varepsilon_{\text{global}}^K,$$

210 *where*

$$\begin{aligned} \varepsilon_{\text{global}}^K &= \min_{\mu} \sum_i \|(I-P)(s^{(i)} - \mu)\|_w^2, \\ \varepsilon_{\text{TEMPO}}^K &= \sum_m \min_{\mu_m} \sum_i \gamma_{im} \|(I-P_m)(s^{(i)} - \mu_m)\|_w^2 \end{aligned}$$

211 *are the global and regime-wise optimal rank- K reconstruction errors. Proof:*
 212 *Appendix A.1.*




Figure 2: Empirical verification of Proposition 1 on 1D Burgers'. *Left*: normalized reconstruction error (log scale); TEMPO lies strictly below the global POD for all $K \geq 1$. *Right*: relative gain $(\varepsilon_{\text{global}}^K - \varepsilon_{\text{TEMPO}}^K) / \varepsilon_{\text{global}}^K$

213 *Adaptive basis update.* Retraining all M NeuralPOD bases at every EM iteration is computationally prohibitive. We introduce a distribution-shift criterion
 214 to determine whether each basis requires updating. Let $\alpha^{(i)}$ denote the global
 215 POD projection of $s^{(i)}$, computed once during initialization and fixed thereafter.
 216 The responsibility-weighted first and second moments of regime m are
 217

$$\mu_m^{(q)} = \frac{1}{N_m^{(q)}} \sum_{i=1}^N \gamma_{im}^{(q)} \alpha^{(i)}, \quad (18)$$

$$\Sigma_m^{(q)} = \frac{1}{N_m^{(q)}} \sum_{i=1}^N \gamma_{im}^{(q)} (\alpha^{(i)} - \mu_m^{(q)})(\alpha^{(i)} - \mu_m^{(q)})^\top. \quad (19)$$

218 The distribution-shift criterion measures the relative change in these moments
 219 between consecutive iterations:

$$\Delta_m^{(q)} = \frac{\|\mu_m^{(q)} - \mu_m^{(q-1)}\|_2}{\|\mu_m^{(q-1)}\|_2 + \epsilon} + \frac{\|\Sigma_m^{(q)} - \Sigma_m^{(q-1)}\|_F}{\|\Sigma_m^{(q-1)}\|_F + \epsilon}. \quad (20)$$

220 where $\epsilon > 0$ is a small numerical stabilizer. $\Delta_m^{(q)}$ measures how much the com-
 221 position of regime m has changed, not just the volume of reassignments. Two
 222 trajectories with very different $\alpha^{(i)}$ swapping regimes produce a large $\Delta_m^{(q)}$ and
 223 trigger a basis update; two similar snapshots doing so do not. Given thresh-
 224 olds $0 < \varepsilon_s < \varepsilon_l$, regime m is SKIPPED if $\Delta_m^{(q)} < \varepsilon_s$, updated INCREMENTALLY
 225 (add or prune one mode) if $\varepsilon_s \leq \Delta_m^{(q)} < \varepsilon_l$, and fully retrained from scratch
 226 (FULL RERUN) otherwise. In the INCREMENTAL case, we compute the residual
 227 of the current basis under the updated responsibilities $\gamma^{(q)}$, add a new mode
 228 if $\sum_i \gamma_{im}^{(q)} \|r_m^{(K_m, i)}\|_w^2 \geq \tau \cdot \sum_i \gamma_{im}^{(q)} \|s^{(i)} - \bar{s}_m^{(q)}\|_w^2$, and prune the last mode if its
 229 weighted contribution falls below ε_p . The FULL RERUN case uses a cold-start
 230 random initialization, as warm-starting may trap gradient descent near the pre-
 231 vious regime's basis.

232 *Convergence and model selection.* The number of regimes M is selected by two
 233 criteria: BIC on the EM log-likelihood (penalizes complexity) and mean assign-
 234 ment entropy H . When they conflict, entropy is decisive: a large gain ΔH
 235 from M to $M+1$ signals a new regime. EM terminates when $\max_m \Delta_m^{(q)} < \varepsilon_c$.
 236 Selected values are $M=3$ (Burgers'), $M=4$ (Navier-Stokes), and $M=5$ (Darcy);
 237 details for Burgers' are in Table A.4 (Appendix A.3). The complete offline pro-
 238 cedure is given in Algorithm 2. Figure 3 shows the resulting regime structure
 239 on 1D Burgers': three manifold branches emerge from unlabelled snapshots, one
 240 per discovered regime, with the partition reflecting joint variability from initial
 241 conditions and viscosity in the solution trajectories.




Figure 3: TEMPO(POD) regime structure on 1D Burgers' ($M=3$, 20 EM iterations). *Left:* UMAP [20] colored by regime. *Center:* same embedding colored by ν . *Right:* GMM fit in POD coefficient space (α_1, α_2) ; stars mark regime centroids μ_m .

242 3.4. TEMPO: Online Gated Operator

243 After Algorithm 2 converges, all regime bases $\{\Phi_m^*\}$ and responsibilities
 244 $\{\gamma_{im}^*\}$ are frozen. The online phase trains a GatingNet $G : \mathbb{R}^{m_s+d_\kappa} \rightarrow \Delta^{M-1}$
 245 that maps (u_0, κ) to regime weights, and M BranchNets $\mathcal{B}_m : \mathbb{R}^{m_s+d_\kappa} \rightarrow \mathbb{R}^{K_m}$
 246 that predict expansion coefficients for each regime. The TEMPO prediction at
 247 query point $y = (x, t) \in \Omega \times \mathcal{T}$ is

$$\hat{s}(y; u_0, \kappa) = \sum_{m=1}^M g_m(u_0, \kappa) \left[\phi_0^{(m)}(y) + \sum_{k=1}^{K_m} b_{m,k}(u_0, \kappa) \phi_k^{(m)}(y) \right], \quad (21)$$

248 where $\mathbf{g} = G(u_0, \kappa)$ and $\mathbf{b}_m = \mathcal{B}_m(u_0, \kappa)$. The BranchNet outputs $\mathbf{b}_m$ predict
 249 the expansion coefficients $\lambda_m^{(i)}$ for new inputs (u_0, κ) not seen in training. The
 250 prediction is mesh-free and evaluable at any $y \in \Omega \times \mathcal{T}$ without retraining.

251 *Training objective.* Writing $g_m^{(i)} = g_m(u_0^{(i)}, \kappa^{(i)})$, the online loss is

$$\begin{aligned}
\mathcal{L}_{\text{online}} = & \underbrace{\frac{1}{N} \sum_{i=1}^N \|s^{(i)} - \hat{s}(\cdot; u_0^{(i)}, \kappa^{(i)})\|_w^2}_{\mathcal{L}_{\text{data}}} \\
& + \underbrace{\eta_{\text{KL}} \frac{1}{N} \sum_{i,m} g_m^{(i)} \log \frac{g_m^{(i)}}{\gamma_{im}^*}}_{\mathcal{L}_{\text{KL}}} \\
& - \underbrace{\eta_H \frac{1}{N} \sum_{i,m} g_m^{(i)} \log g_m^{(i)}}_{-\mathcal{L}_{\text{ent}}}.
\end{aligned} \tag{22}$$

252 $\mathcal{L}_{\text{data}}$ drives reconstruction accuracy. $\mathcal{L}_{\text{KL}}$ anchors the online gating to the offline
253 regime partition: it penalizes deviation of $\mathbf{g}^{(i)}$ from the EM responsibilities γ_i^* ,
254 so the network cannot discard the discovered structure in pursuit of a simpler
255 routing. $\mathcal{L}_{\text{ent}}$ prevents collapse to a single expert, ensuring all M regime bases
256 contribute. The hyperparameters $\eta_{\text{KL}}, \eta_H \geq 0$ trade off these objectives.

257 4. Computational Experiments

258 We evaluate TEMPO on parametric PDE benchmarks, verifying three claims:
259 (i) the offline EM phase recovers interpretable, physically meaningful regime
260 structure from unlabelled snapshots; (ii) locally adapted bases achieve lower re-
261 construction error than a single global basis at equal mode count; (iii) a gated
262 operator with per-regime bases outperforms a single global basis when the data
263 contains structured regimes.

264 4.1. Experimental Setup

265 *Datasets.* We use three parametric PDE benchmarks; the first two are from
266 PDEBench [21].

1D Burgers' equation.

$$u_t + u u_x = \nu u_{xx}, \quad x \in (-1, 1), \quad t \in (0, 2], \tag{23}$$

267 with periodic boundary conditions and random initial conditions. Viscosity
268 $\nu \in \{0.001, 0.01, 0.1, 1.0\}$ spans four qualitatively distinct solution regimes:
269 smooth diffusive profiles at $\nu=1.0$ and near-discontinuous shock structures at
270 $\nu=0.001$. Each value provides $N=9,500$ trajectories on $N_x=1,024$ spatial points
271 and $N_t=201$ time steps (38,000 total).

272 *2D Darcy flow.*

$$-\nabla \cdot (a(x) \nabla u(x)) = \beta, \quad x \in (0, 1)^2, \quad (24)$$

273 with zero Dirichlet boundary conditions. Here $a(x)$ is a spatially varying per-
 274 meability field sampled from a random sum of Gaussians, and β is the constant
 275 source term taking five values: 0.01, 0.1, 1.0, 10.0, 100.0. Each value provides
 276 $N=8,000$ training and 1,000 test samples discretised on a 128×128 uniform grid
 277 ($N_y=16,384$ spatial points), 40,000 samples in total. The operator maps $a(x)$
 278 to the pressure solution $u(x)$. As a steady-state problem, $N_t=1$ and there is no
 time dependence.

2D incompressible Navier-Stokes.

$$\partial_t \omega + \mathbf{u} \cdot \nabla \omega = \frac{1}{\text{Re}} \Delta \omega + f, \quad \mathbf{x} \in (0, 1)^2, \quad t \in (0, 2], \quad (25)$$

279 with periodic boundary conditions, $\omega = \nabla \times \mathbf{u}$, divergence-free $\mathbf{u}$ recovered via
 280 the stream function, and Kolmogorov-type forcing $f(\mathbf{x}) = 0.1(\sin 2\pi(x_1+x_2) +$
 281 $\cos 2\pi(x_1+x_2))$ [13]. Reynolds number $\text{Re} \in \{100, 1000, 3600, 10000\}$ spans
 282 four qualitatively distinct regimes from smooth laminar flow to developed tur-
 283 bulance. Each value provides $N=5,000$ trajectories on a 64×64 periodic grid
 284 (20,000 total; 4,000/1,000 train/test); see Appendix A.2. Operator: $\mathbf{u}_0 \mapsto$
 285 $(\mathbf{u}_t)_{t=1}^{20}, \Delta t=0.1$.

286 *Methods.* We compare DeepONet [4], POD-DeepONet [3], FNO [13], and CNO [14],
 287 all trained jointly on all parameter values via a single global model. We in-
 288 troduce NeuralPOD-DeepONet (learned Fourier NeuralPOD basis with branch
 289 network); TEMPO(POD) and TEMPO(NeuralPOD) (regime-aware bases dis-
 290 covered via EM, routed by gating network). Per-parameter specialists provide
 291 in-distribution reference baselines. All methods implemented by the authors.
 292 Evaluation: mean relative L^2 error (6) per parameter value. Inference times
 293 measured on a single NVIDIA RTX A5000 GPU, per sample.

294 4.2. Main Results

295 *1D Burgers' equation.* Table 1 gives the cross-viscosity error matrix. Per- ν
 296 specialists are unviable across the full range: a $\nu=1.0$ model degrades from
 297 0.025 to 0.461 at $\nu=0.001$. TEMPO(POD) addresses this by training across
 298 all viscosity values, outperforming the joint POD-DeepONet baseline by 13–
 299 50%. TEMPO(NeuralPOD) does not surpass the joint baseline. NeuralPOD
 300 coefficients are harder to predict than POD projections, as seen already in the
 301 single-regime case (0.229 vs. 0.048 at $\nu=0.1$); iterative EM feedback amplifies
 302 the prediction difficulty.

303 *2D Darcy flow.* Table 2 gives the cross- β matrix. Specialists fail out-of-distribution:
 304 Darcy solutions scale linearly with β , driving cross-parameter errors beyond
 305 1000. TEMPO(POD) reduces error by up to $4.75\times$ over the joint POD-DeepONet
 306 baseline ($\beta=0.1$: 0.775 vs. 2.557, $3.3\times$; $\beta=100$: 0.073 vs. 0.347, $4.75\times$; $\beta=1.0$:

2.0 $\times$; $\beta=10$: 1.3 $\times$). TEMPO(NeuralPOD) improves over the global NeuralPOD-DeepONet baseline ($\beta=0.1$: 3.813 vs. 72.9) but does not match TEMPO(POD). FNO and CNO achieve lower absolute error without explicit basis structure; the interpretable regime decomposition that TEMPO provides is not available from these methods.

2D Navier-Stokes. Table 3 gives the cross-Re error matrix. TEMPO(POD) remains within 5–17% of the joint POD-DeepONet baseline but does not improve over it. Unlike Burgers’ and Darcy, solutions at different Re share similar dominant POD directions: the EM assigns near-uniform regime weights, and local bases bring no reconstruction advantage over the global one. TEMPO(NeuralPOD) (κ -initialised, $M=4$) converges but does not close the gap to TEMPO(POD).

Inference speed. TEMPO(POD) evaluates all M branch networks via soft gating, remaining fast relative to FNO and CNO. On Burgers’, TEMPO(POD) requires 0.025 ms per sample versus 9.90 ms for FNO (400 $\times$ faster) and 2.14 ms for CNO (86 $\times$ faster); on Darcy, 0.006 ms versus 0.587 ms (FNO) and 2.24 ms (CNO). The overhead relative to POD-DeepONet is modest (0.025 vs. 0.004 ms on Burgers’), reflecting the additional gating network and M branch evaluations. In latency-critical settings, replacing soft gating with hard routing (selecting a single branch per sample) could reduce this overhead further.

4.3. Phase 1 Reconstruction Advantage

Figure 2 empirically verifies Proposition 1 on 1D Burgers’ ($N=1,500$, $\nu \in \{0.001, 0.1, 1.0\}$) with k -means regime assignments. TEMPO achieves 20–39% lower reconstruction error than the global POD at $K=10$ modes, with the gap growing in both M and K ; no violations of the bound were observed across all $K \in \{1, \dots, 32\}$.

4.4. Regime Discovery

We select $M=3$ based on entropy gain: H increases by +0.30 from $M=2$ to $M=3$ but only +0.18 from $M=3$ to $M=4$, indicating the third regime captures genuine structure that a fourth does not. BIC increases monotonically with M (favoring $M=2$ by complexity alone), so entropy is the decisive criterion here (Table A.4, Appendix A.3). Figure 3 confirms three distinct manifold branches, one per discovered regime. The discovered regimes capture the joint structure of initial conditions and viscosity in the solution space. The per-viscosity gating weights (Table A.5, Appendix A.3) are nearly uniform across $\nu \in \{0.001, 0.1, 1.0\}$, closely matching the global weights $\pi=(0.324, 0.267, 0.409)$: the initial condition determines regime membership, while ν shapes within-regime dynamics. The same initial condition yields a shock at $\nu=0.001$ and a smooth wave at $\nu=1.0$. At inference, ν is passed to the branch network as a conditioning input, capturing both sources of solution variability.

Table 1: **Relative L^2 test error on 1D Burgers' equation.** Each per- ν specialist is trained on a single viscosity and evaluated on all three jointly-trained values; underlined entries are in-distribution. The large off-diagonal errors motivate joint training with regime discovery. All jointly-trained methods use $\nu \in \{0.001, 0.1, 1.0\}$. **Bold**: best per column among jointly-trained methods.

Method	Train ν	$\nu=0.001$	$\nu=0.1$	$\nu=1.0$	Infer. (ms)
<i>Per-ν specialists (single-viscosity training)</i>					
DeepONet [4]	0.001	<u>0.536</u>	0.474	0.382	0.037
	0.1	0.535	<u>0.471</u>	0.349	
	1.0	0.547	0.483	<u>0.283</u>	
POD-DeepONet [3]	0.001	<u>0.273</u>	0.297	0.639	0.004
	0.01	0.274	0.274	0.614	
	0.1	0.349	<u>0.048</u>	0.344	
	1.0	0.461	0.260	<u>0.025</u>	
NeuralPOD-DeepONet (<i>ours</i>)	0.001	<u>0.392</u>	0.288	0.535	3.53
	0.01	0.391	0.284	0.541	
	0.1	0.426	<u>0.229</u>	0.371	
	1.0	0.483	0.315	<u>0.183</u>	
FNO [13]	0.001	<u>0.426</u>	0.529	0.623	9.90
	0.01	0.454	0.401	0.492	
	0.1	0.650	<u>0.304</u>	0.521	
	1.0	1.161	<u>0.919</u>	<u>0.158</u>	
CNO [14] (2.0M)	0.01	0.687	0.955	2.344	2.14
	0.1	0.983	<u>0.199</u>	0.471	
	1.0	1.870	1.084	<u>0.125</u>	
<i>Joint training ($\nu \in \{0.001, 0.1, 1.0\}$)</i>					
DeepONet [4]	all	0.533	0.472	0.289	0.037
POD-DeepONet [3]	all	0.460	0.339	0.210	0.004
FNO [13]	all	0.427	0.298	0.202	9.90
CNO [14] (2.0M)	all	0.361	0.162	0.133	2.14
TEMPO(POD) (<i>ours</i>)	$M=2$	0.319	0.150	0.159	0.025
	$M=3$	0.317	0.168	0.182	
TEMPO(NeuralPOD) (<i>ours</i>)	$M=2$	0.536	0.458	0.326	0.024
	$M=3$	0.542	0.456	0.288	

Table 2: **Relative L^2 test error on 2D Darcy flow.** Each per- β specialist is trained on a single source value and evaluated on all five; underlined entries are in-distribution. Off-diagonal errors exceed 1 because the Darcy solution scales with β , so a specialist trained on $\beta=100$ predicts amplitudes $\sim 10^4$ times larger than the ground truth at $\beta=0.01$. TEMPO trains jointly on $\beta \in \{0.1, 1, 10, 100\}$; the $\beta=0.01$ regime is excluded from joint training (—). **Bold**: best per column among jointly-trained methods.

Method	Train β	$\beta=0.01$	$\beta=0.1$	$\beta=1.0$	$\beta=10$	$\beta=100$	Infer. (ms)
<i>Per-β specialists (single-value training)</i>							
DeepONet [4]	0.01	<u>4.197</u>	0.615	0.944	0.994	0.999	0.025
	0.1	>10	<u>0.851</u>	0.875	0.987	0.999	
	1.0	>10	>1	<u>0.433</u>	0.902	0.972	
	10	>100	>10	>1	<u>0.243</u>	0.705	
	100	>1000	>100	>10	>1	<u>0.078</u>	
POD-DeepONet [3]	0.01	<u>0.449</u>	0.791	0.961	0.996	>1	0.001
	0.1	>1	<u>0.148</u>	0.871	0.986	0.999	
	1.0	>10	>1	<u>0.048</u>	0.897	0.990	
	10	>100	>10	>1	<u>0.035</u>	0.900	
	100	>1000	>100	>10	>1	<u>0.042</u>	
NeuralPOD-DeepONet (<i>ours</i>)	0.01	<u>1.407</u>	0.786	0.962	0.996	>1	0.168
	0.1	>1	<u>0.421</u>	0.902	0.989	0.999	
	1.0	>10	>1	<u>0.135</u>	0.901	0.990	
	10	>100	>10	>1	<u>0.050</u>	0.900	
	100	>1000	>100	>10	>1	<u>0.037</u>	
FNO [13]	0.01	<u>0.529</u>	0.791	0.967	0.997	0.9998	0.587
	0.1	6.26	<u>0.144</u>	0.877	0.988	0.999	
	1.0	73.1	8.41	<u>0.034</u>	0.909	0.992	
	10	>100	72.6	8.18	<u>0.013</u>	0.900	
	100	>1000	256.3	43.2	6.13	<u>0.008</u>	
CNO [14] (2.0M)	0.01	<u>0.388</u>	0.929	1.000	1.001	1.000	2.24
	0.1	10.6	<u>0.125</u>	0.948	1.004	1.002	
	1.0	67.0	7.84	<u>0.024</u>	0.896	0.989	
	10	>1000	99.1	8.37	<u>0.009</u>	0.898	
<i>Joint training ($\beta \in \{0.1, 1, 10, 100\}$)</i>							
DeepONet [4]	all	—	76.53	8.503	1.352	0.493	0.025
POD-DeepONet [3]	all	—	2.557	0.363	0.181	0.347	0.001
NeuralPOD-DeepONet (<i>ours</i>)	all	—	72.9	6.92	0.357	0.345	0.168
FNO [13]	all	—	0.380	0.083	0.023	0.008	0.587
CNO [14] (2.0M)	all	—	0.133	0.026	0.008	0.005	2.24
TEMPO(POD) (<i>ours</i>)	$M=5$	—	0.775	0.178	0.136	0.073	0.006
TEMPO(NeuralPOD) (<i>ours</i>)	$M=5$	—	3.813	0.549	0.322	0.311	0.006

Table 3: **Relative L^2 test error on 2D incompressible Navier-Stokes.** Each per-Re specialist is trained on a single Reynolds number and evaluated on all four; underlined entries are in-distribution. Joint methods train on all Re simultaneously. **Bold:** best per column among jointly-trained methods. [†] kappa-initialised variant.

Method	Train Re	Re=100	Re=1000	Re=3600	Re=10000	Infer. (ms)
<i>Per-Re specialists (single-Re training)</i>						
DeepONet [4]	100	<u>1.000</u>	1.002	1.003	1.004	0.135
	1000	1.000	<u>1.000</u>	1.000	1.000	
	3600	1.002	1.000	<u>1.000</u>	1.000	
	10000	1.001	1.000	<u>1.000</u>	<u>1.000</u>	
POD-DeepONet [3]	100	<u>0.082</u>	0.748	0.865	0.877	0.005
	1000	0.667	<u>0.129</u>	0.671	0.812	
	3600	0.660	0.392	<u>0.180</u>	0.623	
	10000	0.934	0.687	0.556	<u>0.345</u>	
NeuralPOD-DeepONet (<i>ours</i>)	1000	0.762	<u>0.198</u>	0.607	0.761	4.87
FNO [13]	100	<u>0.013</u>	0.262	0.368	0.437	0.30
	1000	0.287	<u>0.033</u>	0.138	0.227	
	3600	0.381	0.150	<u>0.039</u>	0.128	
	10000	0.461	0.260	0.126	<u>0.042</u>	
CNO [14]	100	<u>0.014</u>	0.047	0.055	0.058	0.50
	1000	0.027	<u>0.028</u>	0.033	0.035	
	3600	0.030	0.029	<u>0.033</u>	0.035	
	10000	0.030	0.028	0.032	<u>0.034</u>	
<i>Joint training (all Re)</i>						
DeepONet [4]	all	>1	>1	>1	>1	0.135
POD-DeepONet [3]	all	0.092	0.132	0.141	0.145	0.005
FNO [13]	all	0.009	0.019	0.024	0.027	0.30
NeuralPOD-DeepONet (<i>ours</i>)	all	0.700	0.713	0.314	0.320	4.87
CNO [14]	all	0.020	0.035	0.040	0.043	0.50
TEMPO(POD) (<i>ours</i>)	$M=4$	0.097	0.144	0.155	0.170	0.021
TEMPO(NeuralPOD) [†] (<i>ours</i>)	$M=4$	0.595	0.620	0.694	0.627	0.021

346 4.5. Role of the Physical Parameter

347 We ablate the effect of κ at inference by replacing it with zero or a learned pre-
348 diction. On Burgers', TEMPO(NeuralPOD) barely needs κ : removing it costs
349 only $\sim 1\%$ in L2, since u_0 carries sufficient information to determine regime
350 routing on this benchmark. TEMPO(POD) is more sensitive (+37%); how-
351 ever, ν can be predicted from u_0 with 87% accuracy, reducing the gap to +7%.
352 On Darcy, κ is essential for TEMPO(POD): removing β collapses gating to a
353 fixed mixture, increasing error by $\sim 3\times$. TEMPO(NeuralPOD) is less sensitive
354 (+27%), as its nonlinear basis partially compensates. In both cases, β cannot
355 be predicted from $a(x)$, as it is drawn independently in the dataset (classifier
356 accuracy 21%, below chance). Full results are in Table A.6 (Appendix A.3).

357 5. Conclusion

358 Operator learning for parametric PDEs typically relies on a single global
359 representation. When the governing parameter drives qualitative change in
360 solution structure, a single global basis cannot fit all regimes well and prediction
361 quality suffers.

362 TEMPO discovers regime structure from solution trajectories via EM and fits
363 a dedicated basis to each regime. The discovered partition captures variability
364 driven by both initial conditions and the physical parameter. On 1D Burgers',
365 regimes persist across all viscosity values: the same initial condition produces
366 a shock at $\nu=0.001$ and a smooth wave at $\nu=1.0$, yet falls in the same regime
367 throughout, with ν shaping dynamics within each regime. A global model can-
368 not separate these sources of variation; TEMPO recovers their structure from
369 trajectory data without regime-label supervision.

370 TEMPO(POD) outperforms the joint POD-DeepONet baseline by 13–50%
371 on Burgers' and up to $4.75\times$ on Darcy flow, where per-parameter specialists
372 fail to generalize. On Navier-Stokes, solutions at different Re share similar
373 POD directions; TEMPO identifies this and remains within 5–17% of the global
374 baseline. TEMPO(NeuralPOD) extends the framework to nonlinear bases; im-
375 proving nonlinear basis fitting within the EM loop remains an open direction.

376 Several directions remain open. Conditioning basis functions on κ would let
377 each mode encode parameter-dependent dynamics directly. Extending regime
378 discovery to continuous parameter spaces and three-dimensional problems is a
379 natural next step.

References

- [1] J. S. Hesthaven, G. Rozza, B. Stamm, Certified Reduced Basis Methods for Parametrized Partial Differential Equations, SpringerBriefs in Mathematics, Springer, Cham, 2016. doi:10.1007/978-3-319-22470-1.
- [2] N. B. Kovachki, S. Lanthaler, A. M. Stuart, Operator learning: Algorithms and analysis, arXiv (2024). arXiv:2402.15715.
- [3] L. Lu, X. Meng, S. Cai, Z. Mao, S. Goswami, Z. Zhang, G. E. Karniadakis, A comprehensive and fair comparison of two neural operators (with practical extensions) based on FAIR data, Computer Methods in Applied Mechanics and Engineering 393 (2022) 114778. doi:10.1016/j.cma.2022.114778.
- [4] L. Lu, P. Jin, G. Pang, Z. Zhang, G. E. Karniadakis, Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators, Nature Machine Intelligence 3 (2021) 218–229. doi:10.1038/s42256-021-00302-5.
- [5] T. Chen, H. Chen, Universal approximation to nonlinear operators by neural networks with arbitrary activation functions and its application to dynamical systems, IEEE Transactions on Neural Networks 6 (4) (1995) 911–917. doi:10.1109/72.392253.
- [6] J. Reiss, P. Schulze, J. Sesterhenn, V. Mehrmann, The shifted proper orthogonal decomposition: A mode decomposition for multiple transport phenomena, SIAM Journal on Scientific Computing 40 (3) (2018) A1322–A1344. doi:10.1137/17M1140571.
- [7] D. Amsallem, M. J. Zahr, C. Farhat, Nonlinear model order reduction based on local reduced-order bases, International Journal for Numerical Methods in Engineering 92 (10) (2012) 891–916. doi:10.1002/nme.4371.
- [8] M. Hess, A. Alla, A. Quaini, G. Rozza, M. Gunzburger, A localized reduced-order modeling approach for PDEs with bifurcating solutions, Computer Methods in Applied Mechanics and Engineering 351 (2019) 379–403. doi:10.1016/j.cma.2019.03.050.
- [9] T. Daniel, F. Casenave, N. Akkari, D. Ryckelynck, Model order reduction assisted by deep neural networks (ROM-net), Advanced Modeling and Simulation in Engineering Sciences 7 (1) (2020) 16. doi:10.1186/s40323-020-00153-6.
- [10] R. Sharma, V. Shankar, Ensemble and mixture-of-experts deepoanets for operator learning, arXiv (2025). arXiv:2405.11907.
- [11] K. Lee, K. T. Carlberg, Model reduction of dynamical systems on nonlinear manifolds using deep convolutional autoencoders, Journal of Computational Physics 404 (2020) 108973. doi:10.1016/j.jcp.2019.108973.

- [12] C. Mou, B. Lu, G. Lin, Neural-pod: A plug-and-play neural operator framework for infinite-dimensional functional nonlinear proper orthogonal decomposition, arXiv (2026). [arXiv:2602.15632](#).
- [13] Z. Li, N. Kovachki, K. Azizzadenesheli, B. Liu, K. Bhattacharya, A. Stuart, A. Anandkumar, Fourier neural operator for parametric partial differential equations, in: International Conference on Learning Representations, 2021.
- [14] B. Raonić, R. Molinaro, T. De Ryck, T. Rohner, F. Bartolucci, R. Alaifari, S. Mishra, E. de Bézenac, Convolutional neural operators for robust and accurate learning of PDEs, in: Advances in Neural Information Processing Systems, Vol. 36, 2023.
- [15] A. P. Dempster, N. M. Laird, D. B. Rubin, Maximum likelihood from incomplete data via the em algorithm, *Journal of the Royal Statistical Society: Series B (Methodological)* 39 (1) (1977) 1–38.
- [16] G. Berkooz, P. Holmes, J. L. Lumley, The proper orthogonal decomposition in the analysis of turbulent flows, *Annual Review of Fluid Mechanics* 25 (1993) 539–575. doi:10.1146/annurev.fl.25.010193.002543.
- [17] M. Tancik, P. P. Srinivasan, B. Mildenhall, S. Fridovich-Keil, N. Raghavan, U. Singhal, R. Ramamoorthi, J. T. Barron, R. Ng, Fourier features let networks learn high frequency functions in low dimensional domains, in: Advances in Neural Information Processing Systems, Vol. 33, 2020, pp. 7537–7547.
- [18] N. Rahaman, A. Baratin, D. Arpit, F. Draxler, M. Lin, F. A. Hamprecht, Y. Bengio, A. Courville, On the spectral bias of neural networks, in: Proceedings of the 36th International Conference on Machine Learning, Vol. 97 of Proceedings of Machine Learning Research, 2019, pp. 5301–5310.
- [19] R. M. Neal, G. E. Hinton, A view of the EM algorithm that justifies incremental, sparse, and other variants, in: M. I. Jordan (Ed.), *Learning in Graphical Models*, Vol. 89 of NATO ASI Series, Springer, Dordrecht, 1998, pp. 355–368. doi:10.1007/978-94-011-5014-9_12.
- [20] L. McInnes, J. Healy, J. Melville, UMAP: Uniform manifold approximation and projection for dimension reduction, arXiv (2018). [arXiv:1802.03426](#).
- [21] M. Takamoto, T. Praditia, R. Leiteritz, D. MacKinlay, F. Alesiani, D. Pflüger, M. Niepert, PDEBench: An extensive benchmark for scientific machine learning, in: Advances in Neural Information Processing Systems, Vol. 35, 2022.

453 Appendix A.

454 Appendix A.1. Proof of Proposition 1

455 **Lemma 1.** Let $\bar{s} = \frac{1}{N} \sum_i s^{(i)}$ be the global mean and $\bar{s}_m = \frac{1}{N_m} \sum_i \gamma_{im} s^{(i)}$ the
 456 regime- m mean, with $N_m = \sum_i \gamma_{im}$. For any rank- K orthoprojector P and any
 457 regime m ,

$$\sum_i \gamma_{im} \|(I-P)(s^{(i)} - \bar{s}_m)\|_w^2 = \sum_i \gamma_{im} \|(I-P)(s^{(i)} - \bar{s})\|_w^2 - N_m \|(I-P)(\bar{s}_m - \bar{s})\|_w^2.$$

458 *Proof.* Let $c_m = \bar{s}_m - \bar{s}$. Writing $(I-P)(s^{(i)} - \bar{s}_m) = (I-P)(s^{(i)} - \bar{s}) - (I-P)c_m$
 459 and expanding $\|a - b\|_w^2 = \|a\|_w^2 - 2\langle a, b \rangle_w + \|b\|_w^2$, then summing over i with
 460 weights γ_{im} :

$$\begin{aligned} \sum_i \gamma_{im} \|(I-P)(s^{(i)} - \bar{s}_m)\|_w^2 &= \sum_i \gamma_{im} \|(I-P)(s^{(i)} - \bar{s})\|_w^2 \\ &\quad - 2 \left\langle (I-P) \sum_i \gamma_{im} (s^{(i)} - \bar{s}), (I-P)c_m \right\rangle_w \\ &\quad + N_m \|(I-P)c_m\|_w^2. \end{aligned}$$

461 Since $\sum_i \gamma_{im} (s^{(i)} - \bar{s}) = N_m c_m$, the cross-term equals $-2N_m \|(I-P)c_m\|_w^2$; to-
 462 gether with the quadratic term this gives $-N_m \|(I-P)c_m\|_w^2 = -N_m \|(I-P)(\bar{s}_m - \bar{s})\|_w^2$.
 463 $\square$

464 *Proof of Proposition 1.* Let P^* and P_m^* be the globally and regime- m optimal
 465 rank- K projectors. Since P_m^* minimizes the regime- m objective over all rank- K
 466 projectors, substituting the (generally suboptimal) P^* cannot decrease it:

$$\varepsilon_{\text{TEMPO}}^K = \sum_m \sum_i \gamma_{im} \|(I-P_m^*)(s^{(i)} - \bar{s}_m)\|_w^2 \leq \sum_m \sum_i \gamma_{im} \|(I-P^*)(s^{(i)} - \bar{s}_m)\|_w^2.$$

467 Applying Lemma 1 to each m and summing over m , then using $\sum_m \gamma_{im} = 1$ for
 468 each i :

$$\begin{aligned} \sum_m \sum_i \gamma_{im} \|(I-P^*)(s^{(i)} - \bar{s}_m)\|_w^2 &= \sum_i \|(I-P^*)(s^{(i)} - \bar{s})\|_w^2 - \underbrace{\sum_m N_m \|(I-P^*)(\bar{s}_m - \bar{s})\|_w^2}_{\geq 0} \\ &\leq \sum_i \|(I-P^*)(s^{(i)} - \bar{s})\|_w^2 = \varepsilon_{\text{global}}^K. \end{aligned}$$

469 $\square$

470 Appendix A.2. 2D Navier-Stokes Data Generation

471 We use the same equation and forcing as in [13] with a pseudo-spectral solver
 472 in stream-function formulation: diffusion treated implicitly (Crank–Nicolson),
 473 advection explicitly (second-order Adams–Bashforth), and $\frac{2}{3}$ -rule dealiasing on

the 64×64 grid. Simulation time-step $\delta t=0.01$; each recorded snapshot aggregates 10 sub-steps ($\Delta t=0.1$). Initial vorticity fields are random sums of sinusoids,

$$\omega_0(\mathbf{x}) = \sum_{k=1}^7 \frac{a_k}{1+k} [\sin(2\pi k x_1 + \phi_k) + \cos(2\pi k x_2 + \phi_k)], \quad a_k \sim \mathcal{N}(0, 1), \quad \phi_k \sim \mathcal{U}[0, 2\pi].$$

A 5 s burn-in (500 simulation steps) discards transient behavior before recording begins.

Appendix A.3. Additional Experiments

Regime model selection (1D Burgers’). Table A.4 reports BIC, mean assignment entropy H , and mean test error $\bar{\epsilon}$ for $M \in \{2, 3, 4\}$. Table A.5 shows the per-viscosity EM responsibilities alongside global mixture weights π_m .

Table A.4: Regime count selection on 1D Burgers’. BIC = $-2 \log \mathcal{L} + k \log N$ (lower is better; penalizes complexity); H : mean assignment entropy (higher gain signals a new genuine regime); $\bar{\epsilon}$: mean relative L^2 test error. BIC increases monotonically, favoring $M=2$ by complexity alone. **Bold row**: selected configuration ($M=3$), chosen because the entropy gain $\Delta H=0.30$ from $M=2 \rightarrow 3$ substantially exceeds $\Delta H=0.18$ from $M=3 \rightarrow 4$.

M	BIC ($\times 10^4$, $\downarrow$)	H	$\bar{\epsilon}$
2	3.90	0.622	0.209
3	4.92	0.920	0.222
4	5.47	1.099	0.228

Table A.5: Mean EM assignment weights per viscosity (1D Burgers’, TEMPO(POD), $M=3$) vs. global mixture weights π_m . Each row sums to 1.

ν	Regime 1	Regime 2	Regime 3
0.001	0.32	0.28	0.40
0.1	0.32	0.26	0.42
1.0	0.33	0.26	0.41
π_m	0.32	0.27	0.41

Role of the physical parameter (ablation). Table A.6 reports the effect of replacing κ with zero or a learned prediction at inference.

POD vs. NeuralPOD basis comparison (1D Burgers’). Figure A.4 shows the singular-value and residual decay for POD ($K=30$ modes capture 99.99% of the variance) and Fourier NeuralPOD. The corresponding spatiotemporal mode functions (Figure A.5) reveal a structural difference: POD modes are globally supported, while NeuralPOD modes localise adaptively — at $\nu=0.01$ they concentrate near the shock front. Per-sample reconstructions from both bases are compared in Figure A.7; additional NeuralPOD examples are shown in Figure A.6.




Figure A.4: Basis convergence on 1D Burgers' ($\nu=1.0$): POD singular-value decay (*left*); Fourier NeuralPOD residual decay (*right*).




Figure A.5: First K basis modes on 1D Burgers': POD (*top row*) vs. Fourier NeuralPOD (*bottom row*).




Figure A.6: Fourier NeuralPOD representative reconstructions on 1D Burgers' ($\nu=1.0$).




Figure A.7: Phase 1 reconstruction comparison on 1D Burgers' ($\nu=1.0$). Each row shows one sample: ground truth, POD reconstruction, $|\text{POD error}|$, Fourier NeuralPOD reconstruction, $|\text{FNPOD error}|$.




Figure A.8: Space-time reconstructions on 1D Burgers' ($\nu=1.0$).

Table A.6: κ ablation: increase in mean relative L2 error when κ is removed or predicted, relative to the real- κ baseline. For Darcy, β is drawn independently of $a(x)$; prediction from input data is not possible.

	Burgers'		Darcy	
	NPOD ($M=3$)	POD ($M=3$)	NPOD ($M=5$)	POD ($M=5$)
$\kappa = 0$	+0.006 (+1%)	+0.127 (+37%)	+0.334 (+27%)	+0.618 (+219%)
κ predicted from u_0	+0.004 (+1%)	+0.024 (+7%)	not predictable	

493 *Representative reconstructions (1D Burgers')*. Figure A.8 shows space-time predictions for all four methods on 1D Burgers' ($\nu=1.0$). Each panel: ground truth (left), prediction (center), point-wise error (right).

496 *TEMPO regime structure across benchmarks*. Figure A.10 compares the Phase 1 regime structure across all three benchmarks. On Burgers' (NeuralPOD), the partition mirrors the POD variant (Figure 3) with sharper assignments. On Darcy flow, regimes span multiple β values; TEMPO(NeuralPOD) finds a nearly identical partition. On Navier-Stokes, all Re values overlap in POD coefficient space and regimes 2–4 collapse to a few points, explaining why dedicated bases bring no gain.






Figure A.9: TEMPO(POD) EM convergence on 2D Navier-Stokes ($M=4$).

503 *TEMPO reconstructions (2D Darcy flow)*. Figure A.11 shows representative predictions from both TEMPO variants on 2D Darcy flow ($M=5$). Each panel: ground truth (left), prediction (center), point-wise error (right).


506 *TEMPO reconstructions (2D Navier-Stokes)*. Figure A.12 shows representative predictions from TEMPO(POD) and TEMPO(NeuralPOD) on 2D Navier-Stokes ($M=4$); EM convergence is shown in Figure A.9.



(a) 1D Burgers', NeuralPOD ($M=3$, $H=0.23$)



(b) 2D Darcy flow, POD ($M=4$, shown as example)



(c) 2D Navier-Stokes, POD ($M=4$): no Re separation

Figure A.10: TEMPO EM regime structure across all benchmarks. Layout as in Figure 3.







Figure A.11: Space-time reconstructions on 2D Darcy flow ($M=5$).



(a) TEMPO(POD): ground truth (*left*), prediction (*center*), point-wise absolute error (*right*).



(b) TEMPO(POD) — 3D vorticity field.



(c) TEMPO(NeuralPOD) — 3D vorticity field.

Figure A.12: 2D Navier-Stokes ($M=4$): TEMPO(POD) prediction (*top*); TEMPO(POD) and TEMPO(NeuralPOD) 3D vorticity (*middle, bottom*).

Algorithm 1 NEURALBASIS

Require: weighted snapshots $\{s^{(i)}, \gamma_{im}\}_{i=1}^N$, tolerance τ **Ensure:** basis $\phi_0^{(m)}$, $\{\phi_k^{(m)}, \{\lambda_k^{(m,i)}\}_{i=1}^N\}_{k=1}^{K_m}$

- 1: $\bar{s}^{(m)} \leftarrow \frac{1}{N_m} \sum_{i=1}^N \gamma_{im} s^{(i)}$, $N_m = \sum_{i=1}^N \gamma_{im}$ $\triangleright$ responsibility-weighted mean trajectory
- 2: Train $\phi_0^{(m)}$ by minimizing $\sum_j w_j (\phi_0^{(m)}(y_j) - \bar{s}_j^{(m)})^2$
- 3: $r^{(0,i)} \leftarrow s^{(i)} - \phi_0^{(m)}(y)$ for all i ; $k \leftarrow 0$
- 4: $V_m \leftarrow \sum_{i=1}^N \gamma_{im} \|s^{(i)} - \bar{s}^{(m)}\|_w^2$ $\triangleright$ initial unmodeled variance
- 5: **while** $\sum_{i=1}^N \gamma_{im} \|r^{(k,i)}\|_w^2 \geq \tau V_m$ **and** $k < K_{\max}$ **do**
- 6: Jointly minimize over $\phi_{k+1}^{(m)}$ and $\{\lambda_{k+1}^{(m,i)}\}_{i=1}^N$:

$$\sum_{i=1}^N \gamma_{im} \sum_j w_j \left(\lambda_{k+1}^{(m,i)} \phi_{k+1}^{(m)}(y_j) - r_j^{(k,i)} \right)^2$$

- 7: $r^{(k+1,i)} \leftarrow r^{(k,i)} - \lambda_{k+1}^{(m,i)} \phi_{k+1}^{(m)}(y)$ for all i
 - 8: $k \leftarrow k + 1$
 - 9: **end while**
 - 10: **return** $K_m \leftarrow k$, $\phi_0^{(m)}$, $\{\phi_k^{(m)}, \{\lambda_k^{(m,i)}\}_{i=1}^N\}_{k=1}^{K_m}$
-

Algorithm 2 TEMPO

Require: $\mathcal{D} = \{u_0^{(i)}, \kappa^{(i)}, s^{(i)}\}_{i=1}^N$; $M, \varepsilon_s, \varepsilon_l, \varepsilon_p, \varepsilon_c, \tau$

Ensure: Regime bases $\{\Phi_m^*\}$, amplitudes $\{\Lambda_m^*\}$, responsibilities $\{\gamma_{im}^*\}$, weights $\{\pi_m^*\}$

Initialization

- 1: Compute global POD on $\{s^{(i)}\}$; project to $\alpha^{(i)} \in \mathbb{R}^{d_0}$ for all i
- 2: Run GMM-EM on $\{\alpha^{(i)}\} \rightarrow \gamma_{im}^{(0)}, \pi_m^{(0)}, \mu_m^{(0)}, \Sigma_m^{(0)}$
- 3: **for** $m = 1, \dots, M$ **do**
- 4: $\Phi_m^{(0)}, \Lambda_m^{(0)} \leftarrow \text{NEURALBASIS}(\{s^{(i)}, \gamma_{im}^{(0)}\}, \tau)$ ▷ or weighted SVD for TEMPO(POD), see §3.1
- 5: **end for**
- 6: $\sigma^2 \leftarrow (2N)^{-1} \sum_{i,m} \gamma_{im}^{(0)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2$ (14)
- EM iterations*
- 7: **repeat**
- E-step*
- 8: **for** all $i = 1, \dots, N$ and $m = 1, \dots, M$ **do**
- 9: Compute $\hat{s}_m^{(i)}$ via (4) using $\Phi_m^{(q-1)}$ and $\Lambda_m^{(q-1)}$
- 10: $\gamma_{im}^{(q)} \leftarrow \text{Eq. (15)}$
- 11: **end for**
- M-step: mixture weights*
- 12: $\pi_m^{(q)} \leftarrow N_m^{(q)} / N$ for all m (16)
- M-step: bases*
- 13: **for** $m = 1, \dots, M$ **do**
- 14: Compute $\mu_m^{(q)}$ (18), $\Sigma_m^{(q)}$ (19), $\Delta_m^{(q)}$ (20)
- 15: **if** $\Delta_m^{(q)} < \varepsilon_s$ **then**
- 16: $\Phi_m^{(q)} \leftarrow \Phi_m^{(q-1)}$ SKIP
- 17: **else if** $\Delta_m^{(q)} < \varepsilon_l$ **then** INCREMENTAL
- 18: Compute $r_m^{(K_m)}$ under updated $\gamma^{(q)}$; $V_m^{(q)} \leftarrow \sum_i \gamma_{im}^{(q)} \|s^{(i)} - \bar{s}_m^{(q)}\|_w^2$
- 19: **if** $\sum_i \gamma_{im}^{(q)} \|r_m^{(K_m, i)}\|_w^2 \geq \tau \cdot V_m^{(q)}$ **then**
- 20: Train mode $K_m + 1$; $K_m \leftarrow K_m + 1$
- 21: **end if**
- 22: **if** $\|\phi_{K_m}^{(m)}\|_w^2 \cdot \sum_{i=1}^N \gamma_{im}^{(q)} (\lambda_{K_m}^{(m, i)})^2 < \varepsilon_p$ **then**
- 23: $K_m \leftarrow K_m - 1$
- 24: **end if**
- 25: **else** FULL RERUN
- 26: $\Phi_m^{(q)}, \Lambda_m^{(q)} \leftarrow \text{NEURALBASIS}(\{s^{(i)}, \gamma_{im}^{(q)}\}, \tau)$
- 27: **end if**
- 28: **end for**
- 29: **until** $\max_m \Delta_m^{(q)} < \varepsilon_c$
- 30: **return** $\Phi_m^* \leftarrow \Phi_m^{(q)}, \Lambda_m^* \leftarrow \Lambda_m^{(q)}, \gamma_{im}^* \leftarrow \gamma_{im}^{(q)}, \pi_m^* \leftarrow \pi_m^{(q)}$
