
TEMPO: Regime-Aware Operator Learning with Locally Adapted POD Bases

Anonymous Author(s)

Affiliation

Address

email

Abstract

Operator learning methods typically approximate the solution map of a physical system with a single global model—an assumption that fails when solutions change character fundamentally across the parameter space.

We propose TEMPO: a framework that identifies latent dynamical regimes in snapshot data via Expectation-Maximization, builds a dedicated NeuralPOD basis for each, and learns to route new inputs to the right regime at inference time. On multi-regime benchmarks, TEMPO achieves substantially lower errors than POD-DeepONet and vanilla DeepONet, with interpretable structure that reflects the underlying physics.

1 Introduction

Operator learning approximates the solution map of a parametric PDE from data, enabling instant prediction without re-solving the governing equations at each new parameter value [1]. The dominant paradigm—exemplified by DeepONet [2] and POD-DeepONet [3]—trains a single global model over the entire parameter space. This works well when solutions form a simple manifold, but breaks down when the governing parameter drives *qualitatively distinct* dynamics: viscosity separating laminar diffusion from near-discontinuous shocks, or source magnitude producing solutions that differ by orders of magnitude in amplitude. In these multi-regime settings, a single global basis must compromise between incompatible geometries, concentrating error precisely where accurate prediction matters most.

Several methods improve expressivity while retaining a single global representation. POD-PoU [4] trains a partition-of-unity mixture of POD bases with a shared branch, reducing error on sharp-gradient problems but without discovering the underlying regime structure. NeuralPOD [5] constructs nonlinear orthogonal basis functions via greedy residual minimization, representing each mode as a neural network and thus overcoming the linear subspace constraint. Both, however, operate on the full dataset globally and provide no operator for routing new inputs to the appropriate representation at inference.

Clustering snapshots and fitting local bases to each cluster is a classical idea in reduced-order modeling [6, 7]. Daniel et al. [8] extend this to deep learning by clustering solution subspaces on the Grassmann manifold and training a dictionary selector at inference. These methods use hard cluster assignments, operate offline on fixed datasets, and do not couple regime discovery with the operator learning objective—leaving basis quality and assignment accuracy disconnected.

We propose **TEMPO** (*Regime-Aware Operator Learning with Locally Adapted POD Bases*), which closes this gap in a principled probabilistic framework. TEMPO models the trajectory ensemble as a Gaussian mixture and recovers regime structure via Expectation-Maximization (EM): the E-step softly assigns each trajectory to regimes based on reconstruction quality, while the M-step rebuilds

a dedicated NeuralPOD basis per regime so as bases sharpen, assignments improve, and vice versa. After convergence, a GatingNet and per-regime BranchNets are trained jointly, routing new inputs at inference without access to the physical parameter.

Contributions:

- A Gaussian mixture model with EM that discovers M latent regimes and fits a dedicated NeuralPOD basis per regime, requiring no regime labels.
- NeuralPOD-DeepONet: the first deployable operator built on the NeuralPOD basis of [5], pairing it with a branch network that predicts mode coefficients for unseen inputs.
- An adaptive basis update rule based on a distribution-shift criterion, which determines when each regime’s basis warrants retraining and avoids redundant computation across EM iterations.
- On 1D Burgers’ and 2D Darcy flow, a single TEMPO model trained jointly across all parameter values substantially outperforms single-regime baselines, with discovered regimes that reflect intrinsic solution geometry rather than parameter labels.

2 Problem Statement

Let $\Omega \subset \mathbb{R}^d$ be a bounded spatial domain and $\mathcal{T} \subset \mathbb{R}_{\geq 0}$ a time interval.

We are given a dataset of N solution trajectories:

$$\mathcal{D} = \{ (u_0^{(i)}, \kappa^{(i)}, s^{(i)}) \}_{i=1}^N, \quad s^{(i)} = s(\cdot, \cdot; u_0^{(i)}, \kappa^{(i)}) \in L^2(\Omega \times \mathcal{T}), \quad (1)$$

where $u_0^{(i)} \in \mathcal{U} = L^2(\Omega)$ is the initial condition and $\kappa^{(i)} \in \mathcal{K}$ the physical parameter. Each trajectory $s^{(i)}$ is discretized on $N_y = N_t N_x$ quadrature points $\{y_j = (x_j, t_j), w_j\}_{j=1}^{N_y}$ covering the product grid $\Omega \times \mathcal{T}$, with weighted norm $\|f\|_w^2 = \sum_j w_j f(y_j)^2$.

Definition 1 (Solution operator). *The solution operator*

$$\mathcal{G} : \mathcal{U} \times \mathcal{K} \rightarrow L^2(\Omega \times \mathcal{T})$$

maps each pair (u_0, κ) to the full space-time solution: $\mathcal{G}(u_0, \kappa) = s(\cdot, \cdot; u_0, \kappa)$.

The operator $\mathcal{G}$ is unknown and expensive to evaluate directly. The *operator learning problem* is to construct $\hat{\mathcal{G}} \approx \mathcal{G}$ from $\mathcal{D}$ alone, generalizing to unseen $(u_0, \kappa) \notin \mathcal{D}$ without solving the PDE.

Multi-regime assumption. When solutions change character fundamentally across $\mathcal{K}$, the optimal low-dimensional representation differs from regime to regime, and no single global basis captures the full solution diversity. We formalize this observation via a latent variable model over the trajectory ensemble.

Definition 2 (Generative model). *Each trajectory arises from one of M latent regimes. The regime label $z_i \in \{1, \dots, M\}$ is unobserved and drawn from a categorical prior,*

$$z_i \sim \text{Categorical}(\pi_1, \dots, \pi_M), \quad \pi_m \geq 0, \quad \sum_m \pi_m = 1. \quad (2)$$

Conditioned on $z_i = m$, the trajectory is distributed as

$$(s^{(i)} \mid z_i = m) \sim \mathcal{N}(\hat{s}_m^{(i)}, \sigma^2 W^{-1}), \quad (3)$$

where $W = \text{diag}(w_1, \dots, w_{N_y})$ is the quadrature weight matrix, $\sigma^2 > 0$ controls the softness of regime assignments, and $\hat{s}_m^{(i)}$ is the regime- m approximation to $s^{(i)}$.

The generative model leaves $\hat{s}_m^{(i)}$ unspecified. We require only that it admits a *low-rank trajectory decomposition*:

$$\hat{s}_m^{(i)}(y) = \phi_0^{(m)}(y) + \sum_{k=1}^{K_m} \lambda_k^{(m,i)} \phi_k^{(m)}(y), \quad y = (x, t) \in \Omega \times \mathcal{T}, \quad (4)$$

where $\phi_k^{(m)} : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$ are spatiotemporal mode functions and $\lambda_k^{(m,i)} \in \mathbb{R}$ a scalar amplitude for trajectory i on mode k . This form encompasses POD [3] and learned neural bases [5]; in this work we adopt the latter.

Definition 3 (Regime approximator). The parameters of regime m are $\Phi_m = \{\phi_k^{(m)}\}_{k=0}^{K_m}$ and scalar amplitudes $\Lambda_m = \{\lambda_k^{(m,i)}\}_{k=1, i=1}^{K_m, N}$, where $\phi_0^{(m)} : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$ is the regime mean field, $\{\phi_k^{(m)}\}_{k \geq 1}$ spatiotemporal mode functions, $\lambda_k^{(m,i)} \in \mathbb{R}$ the scalar amplitude of trajectory i on mode k , and $K_m \leq K_{\max}$ the adaptive mode count.

The learning problem decomposes into two phases: discovering the latent regime structure from $\mathcal{D}$ *offline*, then training a deployable operator for fast inference *online*.

Phase 1: Offline regime discovery. Since the regime labels z_i in (3) are unobserved, we recover the bases $\{\Phi_m\}$, mixture weights π , and noise level σ^2 by maximizing the observed-data log-likelihood:

$$(\Phi^*, \pi^*, \sigma^{2*}) = \arg \max_{\Phi, \pi, \sigma^2} \sum_{i=1}^N \log \sum_{m=1}^M \frac{\pi_m \det(W)^{1/2}}{(2\pi\sigma^2)^{N_y/2}} \exp\left(-\frac{\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2}{2\sigma^2}\right). \quad (5)$$

The solution yields soft regime responsibilities $\gamma_{im}^* = p(z_i=m \mid s^{(i)}, \Phi^*, \pi^*, \sigma^{2*})$ that serve as supervision for Phase 2.

Phase 2: Online operator learning. With bases $\{\Phi_m^*\}$ and responsibilities $\{\gamma_{im}^*\}$ frozen, we train a gating network (*GatingNet*) and M branch networks (*BranchNets*) that map a new input (u_0, κ) to a weighted combination of local regime predictions. The precise architecture and loss are given in Section 3.4. The quality of the final operator is measured by the mean relative L^2 error:

$$\varepsilon = \frac{1}{N_{\text{test}}} \sum_{i=1}^{N_{\text{test}}} \frac{\|s^{(i)} - \hat{s}(\cdot; u_0^{(i)}, \kappa^{(i)})\|_w}{\|s^{(i)}\|_w}. \quad (6)$$

3 Method

All methods studied in this work share a common two-phase structure: an *offline* basis-fitting phase that constructs a low-rank approximation to the solution manifold from snapshots in $\mathcal{D}$, and an *online* operator-learning phase that trains a branch network to predict expansion coefficients from a new input (u_0, κ) . The four methods differ along two independent axes:

- **Basis type.** A *linear* (POD) basis uses SVD of the (weighted) snapshot matrix; a *nonlinear* (NeuralPOD) basis learns spatiotemporal mode networks by sequential residual minimization.
- **Number of regimes.** A *single-regime* operator fits one global basis to all snapshots; a *multi-regime* operator (TEMPO) discovers M latent regimes via EM and assigns a dedicated basis to each.

This gives four combinations: POD-DeepONet [3] (linear, global), NeuralPOD-DeepONet (nonlinear, global, our extension of [5]), TEMPO(POD) (linear, M regimes), and TEMPO(NeuralPOD) (nonlinear, M regimes).

3.1 Basis Representations

3.1.1 POD Basis

Let $\phi_0 := \bar{s} = \frac{1}{N} \sum_i s^{(i)}$ be the mean field and $\tilde{S} \in \mathbb{R}^{N_y \times N}$ the matrix with columns $\tilde{s}^{(i)} = s^{(i)} - \phi_0$. The rank- K truncated SVD $\tilde{S} \approx U_K \Sigma_K V_K^\top$ yields orthonormal modes $\{\phi_k\}_{k=1}^K$ (columns of U_K) and coefficients $\lambda_k^{(i)} = [\Sigma_K V_K^\top]_{ki}$; the approximation

$$\hat{s}^{(i)} = \phi_0 + \sum_{k=1}^K \lambda_k^{(i)} \phi_k \quad (7)$$

107 is optimal in the L^2 sense among all rank- K affine representations.

108 For TEMPO(POD), the M-step (17) is solved analytically: set $\phi_0^{(m)} := \bar{s}_m = \sum_i \gamma_{im} s^{(i)}$ and apply
 109 SVD to the weighted centered matrix with columns $\sqrt{\gamma_{im}}(s^{(i)} - \phi_0^{(m)})$. This is the closed-form
 110 linear analogue of Algorithm 1.

111 3.1.2 NeuralPOD Basis

112 To overcome the linearity constraint, we parametrize each mode function $\phi_k : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$ as a
 113 Fourier feature network [9]; we call the resulting basis a *Fourier NeuralPOD* basis:

$$\phi_k(y) = f_{\theta_k}(\psi(y)), \quad \psi(y) = [\sin(2\pi B y), \cos(2\pi B y)], \quad (8)$$

114 where $B \in \mathbb{R}^{F \times d}$ is a fixed random matrix drawn once at initialization and $f_{\theta_k} : \mathbb{R}^{2F} \rightarrow \mathbb{R}$ is
 115 a small MLP with Tanh activations. Plain MLPs exhibit spectral bias [10]; the Fourier encoding
 116 resolves this. For multi-scale coverage, the rows of B are partitioned evenly across L frequency
 117 scales $\sigma_1 < \dots < \sigma_L$: the ℓ -th block is drawn from $\mathcal{N}(0, \sigma_\ell^2 I)$. The mean field ϕ_0 uses the same
 118 architecture.

119 The scalar amplitudes $\lambda_k^{(i)} \in \mathbb{R}$ are direct learnable parameters, one per training trajectory, optimized
 120 jointly with ϕ_k . After Phase 1 they are frozen and serve as regression targets for the BranchNet in
 121 Phase 2.

122 Modes are learned by *greedy residual minimization* [5]. Let $r^{(0,i)} = s^{(i)} - \phi_0$, where ϕ_0 is fitted to
 123 the snapshot mean $\bar{s}$ defined above. For each subsequent mode $k \geq 1$, given the residual

$$r^{(k-1,i)}(y) = s^{(i)}(y) - \phi_0(y) - \sum_{\ell=1}^{k-1} \lambda_\ell^{(i)} \phi_\ell(y), \quad (9)$$

124 we solve

$$\min_{\phi_k, \{\lambda_k^{(i)}\}} \sum_{i=1}^N \sum_j w_j \left(\lambda_k^{(i)} \phi_k(y_j) - r_j^{(k-1,i)} \right)^2 \quad (10)$$

125 (the global case with uniform weights; the γ -weighted generalisation used in TEMPO is Algo-
 126 rithm 1). Mode addition stops when the residual falls below fraction τ of the initial variance,

$$\sum_{i=1}^N \|r^{(k,i)}\|_w^2 < \tau \cdot \sum_{i=1}^N \|s^{(i)} - \bar{s}\|_w^2, \quad (11)$$

127 or $k = K_{\max}$. Each mode captures the largest remaining variance component the nonlinear ana-
 128 logue of Gram-Schmidt orthogonalization.

129 3.2 Single-Regime Operators

130 3.2.1 POD-DeepONet

131 POD-DeepONet [3] replaces the learned trunk network of DeepONet [2] with the fixed POD basis.
 132 After computing the global modes $\{\phi_k\}_{k=1}^K$ and training coefficients $\{\lambda^{(i)}\}$ offline, a branch network
 133 $\mathcal{B}_\theta : \mathbb{R}^{m_s + d_\kappa} \rightarrow \mathbb{R}^K$ is trained online by regression against these coefficients:

$$\mathcal{L}_{\text{branch}} = \frac{1}{N} \sum_{i=1}^N \|\mathcal{B}_\theta(u_0^{(i)}, \kappa^{(i)}) - \lambda^{(i)}\|^2. \quad (12)$$

134 The prediction at any $(x, t) \in \Omega \times \mathcal{T}$ is

$$\hat{s}(x, t; u_0, \kappa) = \phi_0(x, t) + \sum_{k=1}^K \mathcal{B}_{\theta,k}(u_0, \kappa) \phi_k(x, t). \quad (13)$$

3.2.2 NeuralPOD-DeepONet

Mou et al. [5] introduce NeuralPOD as a dimensionality reduction for PDE solutions, establishing the greedy basis-fitting procedure of Section 3.1.2. However, they do not construct a deployable solution operator: no network is proposed to predict the mode coefficients for an unseen initial condition. We close this gap by substituting the NeuralPOD basis into the POD-DeepONet framework; we call this combination *NeuralPOD-DeepONet*.

The procedure follows POD-DeepONet exactly with the Fourier NeuralPOD basis substituted for POD. Offline fitting yields coefficients $\lambda^{(i)} \in \mathbb{R}^K$; online, a branch network minimizes (12) and prediction follows (13). At equal mode count K , the nonlinear basis spans a strictly richer function space than POD.

3.3 TEMPO: Offline Regime Discovery

A single global basis whether linear or nonlinear must compromise between qualitatively distinct regimes, as the optimal basis geometry differs fundamentally between them. TEMPO avoids this by constructing a dedicated basis for each of M latent regimes, discovered from unlabelled trajectories via EM on the generative model of Definition 2.

Initialization. A global POD is computed on all N trajectories; each trajectory is projected to a P -dimensional coefficient vector $\alpha^{(i)} \in \mathbb{R}^P$. Running a standard GMM on $\{\alpha^{(i)}\}$ yields initial responsibilities $\gamma_{im}^{(0)}$, weights $\pi_m^{(0)}$, and per-regime statistics $(\mu_m^{(0)}, \Sigma_m^{(0)})$. Each regime basis $\Phi_m^{(0)}$ is then fitted to the responsibility-weighted ensemble: for TEMPO(NeuralPOD) via WEIGHTEDNEURALPOD (Algorithm 1); for TEMPO(POD) via the weighted SVD of Section 3.1. The noise level σ^2 is calibrated from the initial reconstructions rather than treated as a free hyperparameter:

$$\sigma^2 \leftarrow \frac{1}{2N} \sum_{i=1}^N \sum_{m=1}^M \gamma_{im}^{(0)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2, \quad (14)$$

which ties the softness of the E-step directly to the quality of the initial bases.

E-step. At iteration q , the responsibility of trajectory i for regime m is updated by Bayes' rule on the Gaussian likelihood (3):

$$\gamma_{im}^{(q)} = \frac{\pi_m^{(q-1)} \exp\left(-\frac{\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2}{2\sigma^2}\right)}{\sum_{j=1}^M \pi_j^{(q-1)} \exp\left(-\frac{\|s^{(i)} - \hat{s}_j^{(i)}\|_w^2}{2\sigma^2}\right)}, \quad (15)$$

where $\hat{s}_m^{(i)}$ is evaluated via (4) using $\Phi_m^{(q-1)}$ and $\Lambda_m^{(q-1)}$. As the bases improve in the M-step, the reconstructions $\hat{s}_m^{(i)}$ sharpen and the E-step automatically refines regime assignments in response distinguishing TEMPO from a GMM applied to a fixed feature space. When κ drives large amplitude variation (as with $\beta \in [0.01, 100]$ in the Darcy experiments), replacing the squared norm with the relative squared norm $\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2 / \|s^{(i)}\|_w^2$ in both (15) and (14) gives a scale-invariant analogue. When κ is observed, the GMM initialization can operate in $\log \kappa$ space, and an optional κ -prior term in the E-step anchors assignments to the known parameter structure.

M-step. With responsibilities fixed, the M-step gives a closed-form update for the mixture weights,

$$\pi_m^{(q)} \leftarrow \frac{N_m^{(q)}}{N}, \quad N_m^{(q)} = \sum_{i=1}^N \gamma_{im}^{(q)}, \quad (16)$$

and a weighted reconstruction problem for each basis:

$$\min_{\Phi_m, \Lambda_m} \sum_{i=1}^N \gamma_{im}^{(q)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2, \quad (17)$$

solved by weighted SVD for TEMPO(POD) or WEIGHTEDNEURALPOD (Algorithm 1) for TEMPO(NeuralPOD). Since gradient descent solves the NeuralPOD subproblem without guaranteeing global optimality, TEMPO is a Generalized EM algorithm [11], which guarantees that the observed-data log-likelihood is non-decreasing at every iteration.

172 **Adaptive basis update.** Retraining all M NeuralPOD bases at every EM iteration is computationally prohibitive. We introduce a distribution-shift criterion to determine whether each basis requires
 173 updating. Let $\alpha^{(i)}$ denote the global POD projection of $s^{(i)}$, computed once during initialization and
 174 fixed thereafter. The responsibility-weighted first and second moments of regime m are
 175

$$\mu_m^{(q)} = \frac{1}{N_m^{(q)}} \sum_{i=1}^N \gamma_{im}^{(q)} \alpha^{(i)}, \quad (18)$$

$$\Sigma_m^{(q)} = \frac{1}{N_m^{(q)}} \sum_{i=1}^N \gamma_{im}^{(q)} (\alpha^{(i)} - \mu_m^{(q)})(\alpha^{(i)} - \mu_m^{(q)})^\top. \quad (19)$$

176 The distribution-shift criterion measures the relative change in these moments between consecutive
 177 iterations:

$$\Delta_m^{(q)} = \frac{\|\mu_m^{(q)} - \mu_m^{(q-1)}\|_2}{\|\mu_m^{(q-1)}\|_2 + \epsilon} + \frac{\|\Sigma_m^{(q)} - \Sigma_m^{(q-1)}\|_F}{\|\Sigma_m^{(q-1)}\|_F + \epsilon}. \quad (20)$$

178 where $\epsilon > 0$ is a small numerical stabiliser. $\Delta_m^{(q)}$ measures how much the composition of regime
 179 m has changed, not just the volume of reassignments. Two trajectories with very different $\alpha^{(i)}$
 180 swapping regimes produce a large $\Delta_m^{(q)}$ and trigger a basis update; two similar snapshots doing so
 181 do not. Given thresholds $0 < \varepsilon_s < \varepsilon_l$, the update decision at iteration q follows a three-way rule:

	Condition	Action	Rationale
182	$\Delta_m^{(q)} < \varepsilon_s$	SKIP	Data manifold unchanged; basis remains accurate
	$\varepsilon_s \leq \Delta_m^{(q)} < \varepsilon_l$	INCREMENTAL	Moderate drift; add or prune one mode
	$\Delta_m^{(q)} \geq \varepsilon_l$	FULL RERUN	Large shift; rebuild from random initialization

183 In the INCREMENTAL case, we compute the residual of the current basis under the updated responsi-
 184 bilities $\gamma^{(q)}$, add a new mode if $\sum_i \gamma_{im}^{(q)} \|r_m^{(K_m, i)}\|_w^2 \geq \tau \cdot \sum_i \gamma_{im}^{(q)} \|s^{(i)} - \bar{s}_m^{(q)}\|_w^2$, and prune the last
 185 mode if its weighted contribution falls below ε_p . The FULL RERUN case uses a cold-start random
 186 initialization, as warm-starting may trap gradient descent in the geometry of the previous regime.

187 **Convergence and model selection.** EM terminates when $\max_m \Delta_m^{(q)} < \varepsilon_c$. The number of
 188 regimes $M \in \{2, 3, 4\}$ is selected by BIC on the GMM fit, with mean assignment entropy as a
 189 secondary criterion; see Table 3. The complete offline procedure is given in Algorithm 2.

190 3.4 TEMPO: Online Gated Operator

191 After Algorithm 2 converges, all regime bases $\{\Phi_m^*\}$ and responsibilities $\{\gamma_{im}^*\}$ are frozen. The
 192 online phase trains a GatingNet $\mathcal{W} : \mathbb{R}^{m_s+d_\kappa} \rightarrow \Delta^{M-1}$ that maps (u_0, κ) to regime weights, and
 193 M BranchNets $\mathcal{B}_m : \mathbb{R}^{m_s+d_\kappa} \rightarrow \mathbb{R}^{K_m}$ that predict expansion coefficients for each regime. The
 194 TEMPO prediction at query point $y = (x, t) \in \Omega \times \mathcal{T}$ is

$$\hat{s}(y; u_0, \kappa) = \sum_{m=1}^M g_m(u_0, \kappa) \left[\phi_0^{(m)}(y) + \sum_{k=1}^{K_m} b_{m,k}(u_0, \kappa) \phi_k^{(m)}(y) \right], \quad (21)$$

195 where $\mathbf{g} = \mathcal{W}(u_0, \kappa)$ and $\mathbf{b}_m = \mathcal{B}_m(u_0, \kappa)$. The BranchNet outputs $\mathbf{b}_m$ predict the expansion co-
 196 efficients $\lambda_m^{(i)}$ for new inputs (u_0, κ) not seen in training. The prediction is mesh-free and evaluable
 197 at any $y \in \Omega \times \mathcal{T}$ without retraining.

198 **Training objective.** Writing $g_m^{(i)} = g_m(u_0^{(i)}, \kappa^{(i)})$, the online loss is

$$\mathcal{L}_{\text{online}} = \underbrace{\frac{1}{N} \sum_{i=1}^N \|s^{(i)} - \hat{s}(\cdot; u_0^{(i)}, \kappa^{(i)})\|_w^2}_{\mathcal{L}_{\text{data}}} + \underbrace{\lambda_{\text{KL}} \frac{1}{N} \sum_{i,m} g_m^{(i)} \log \frac{g_m^{(i)}}{\gamma_{im}^*}}_{\mathcal{L}_{\text{KL}}} - \underbrace{\lambda_H \frac{1}{N} \sum_{i,m} g_m^{(i)} \log g_m^{(i)}}_{-\mathcal{L}_{\text{ent}}}. \quad (22)$$

$\mathcal{L}_{\text{data}}$ drives reconstruction accuracy. $\mathcal{L}_{\text{KL}}$ anchors the online gating to the offline regime partition: it penalizes deviation of $\mathbf{g}^{(i)}$ from the EM responsibilities γ_i^* , so the network cannot discard the discovered structure in pursuit of a simpler routing. $\mathcal{L}_{\text{ent}}$ prevents collapse to a single expert, ensuring all M regime bases contribute. The hyperparameters $\lambda_{\text{KL}}, \lambda_H \geq 0$ trade off these objectives.

4 Computational Experiments

We evaluate TEMPO on parametric PDE benchmarks, verifying three claims: (i) the offline EM phase recovers interpretable, physics-consistent regime structure from unlabelled snapshots; (ii) locally adapted bases achieve lower reconstruction error than a single global basis at equal mode count; (iii) a jointly-trained gated operator with regime-specific bases outperforms a single jointly-trained basis without regime discovery.

4.1 Experimental Setup

Datasets. We use two parametric PDE benchmarks from PDEBench [12].

1D Burgers’ equation.

$$u_t + u u_x = \nu u_{xx}, \quad x \in (-1, 1), \quad t \in (0, 2], \quad (23)$$

with periodic boundary conditions and random initial conditions. Viscosity $\nu \in \{0.001, 0.01, 0.1, 1.0\}$ spans four qualitatively distinct solution regimes: smooth diffusive profiles at $\nu=1.0$ and near-discontinuous shock structures at $\nu=0.001$. Each value provides $N=9,500$ trajectories on $N_x=1,024$ spatial points and $N_t=201$ time steps (38,000 total).

2D Darcy flow.

$$-\nabla \cdot (a(x) \nabla u(x)) = \beta, \quad x \in (0, 1)^2, \quad (24)$$

with zero Dirichlet boundary conditions. Here $a(x)$ is a spatially varying permeability field sampled from a random sum of Gaussians, and $\beta \in \{0.01, 0.1, 1.0, 10.0, 100.0\}$ is the constant source term. Each value provides $N=8,000$ training and 1,000 test samples discretised on a 128×128 uniform grid ($N_y=16,384$ spatial points), 40,000 trajectories in total. The operator maps $a(x)$ to the pressure solution $u(x)$.

Methods. All methods are implemented and trained by us. POD-DeepONet [3] and NeuralPOD-DeepONet (*ours*) are trained one model per ν (or β) value and evaluated in-distribution; they serve as a reference ceiling for jointly-trained approaches. POD-DeepONet trained jointly on all parameter values without regime discovery is the primary baseline for TEMPO. TEMPO(POD) uses responsibility-weighted SVD per regime; TEMPO(NeuralPOD) uses Fourier NeuralPOD bases. Both TEMPO variants train a single model on all parameter values jointly, without access to ν or β at inference.

Metric. All methods are evaluated by the mean relative L^2 error ε defined in (6), computed per parameter value and reported in Tables 1–2.

4.2 Main Results

1D Burgers’ equation. Table 1 reports the full cross-viscosity error matrix for the Burgers’ benchmark. Each per- ν specialist achieves low error in-distribution (diagonal entries), but degrades sharply outside its training regime: for example, a model trained on $\nu=0.1$ reaches 0.048 in-distribution yet 0.461 on $\nu=0.001$, a ten-fold increase. This cross-regime failure motivates joint training with automatic regime discovery. TEMPO is compared against a jointly-trained POD-DeepONet without regime discovery; TEMPO improves on this baseline by learning regime-specific bases.

2D Darcy flow. Table 2 shows the cross- β error matrix for Darcy flow. Per- β specialists fail even more severely than in Burgers: the Darcy solution scales linearly with β , so a model trained on one value predicts completely wrong amplitudes on others. Errors grow by roughly one order of magnitude per decade in β , reaching >1000 for the most distant pair. TEMPO trains jointly on $\beta \in \{0.1, 1, 10, 100\}$ and recovers reasonable accuracy across all regimes.

Table 1: **Relative L^2 test error on 1D Burgers’ equation.** Each per- ν specialist is trained on a single viscosity and evaluated on all four; underlined entries are in-distribution. The large off-diagonal errors motivate joint training with regime discovery. TEMPO trains jointly on $\nu \in \{0.001, 0.1, 1.0\}$; the $\nu=0.01$ column is omitted for jointly-trained methods (—). **Bold:** best per column among jointly-trained methods.

Method	Train ν	$\nu=0.001$	$\nu=0.01$	$\nu=0.1$	$\nu=1.0$
<i>Per-ν specialists (single-viscosity training)</i>					
POD-DeepONet [3]	0.001	<u>0.273</u>	0.215	0.297	0.639
	0.01	0.274	<u>0.214</u>	0.274	0.614
	0.1	0.349	0.299	<u>0.048</u>	0.344
	1.0	0.461	0.425	<u>0.260</u>	<u>0.025</u>
NeuralPOD-DeepONet (<i>ours</i>)	0.001	<u>0.392</u>	0.352	0.288	0.535
	0.01	0.391	<u>0.351</u>	0.284	0.541
	0.1	0.426	0.388	<u>0.229</u>	0.371
	1.0	0.483	0.451	0.315	<u>0.183</u>
<i>Joint training ($\nu \in \{0.001, 0.1, 1.0\}$)</i>					
POD-DeepONet [3]	all	—	—	—	—
TEMPO(POD) (<i>ours</i>)	$M=3$	0.317	—	0.168	0.182
TEMPO(NeuralPOD) (<i>ours</i>)	$M=3$	0.571	—	0.502	0.376

Table 2: **Relative L^2 test error on 2D Darcy flow.** Each per- β specialist is trained on a single source value and evaluated on all five; underlined entries are in-distribution. Off-diagonal errors exceed 1 because the Darcy solution scales with β , so a specialist trained on $\beta=100$ predicts amplitudes $\sim 10^4$ times larger than the ground truth at $\beta=0.01$. TEMPO trains jointly on $\beta \in \{0.1, 1, 10, 100\}$; the $\beta=0.01$ regime is excluded from joint training (—). **Bold:** best per column among jointly-trained methods.

Method	Train β	$\beta=0.01$	$\beta=0.1$	$\beta=1.0$	$\beta=10$	$\beta=100$
<i>Per-β specialists (single-value training)</i>						
POD-DeepONet [3]	0.01	<u>0.449</u>	0.791	0.961	0.996	>1
	0.1	>1	<u>0.148</u>	0.871	0.986	0.999
	1.0	>10	>1	<u>0.048</u>	0.897	0.990
	10	>100	>10	>1	<u>0.035</u>	0.900
	100	>1000	>100	>10	>1	<u>0.042</u>
<i>Joint training ($\beta \in \{0.1, 1, 10, 100\}$)</i>						
POD-DeepONet [3]	all	—	—	—	—	—
TEMPO(POD) (<i>ours</i>)	$M=5$	—	0.775	0.178	0.136	0.073
TEMPO(NeuralPOD) (<i>ours</i>)	$M=5$	—	3.813	0.549	0.322	0.311

244 4.3 Regime Discovery

245 **Number of regimes.** We sweep $M \in \{2, 3, 4\}$ and evaluate the EM reconstruction criterion (BIC,
246 higher indicates better mixture fit), mean assignment entropy $H = -\frac{1}{N} \sum_{i,m} \gamma_{im} \log \gamma_{im}$, and
247 overall test error $\bar{\varepsilon}$. Table 3 reports all three. $M=3$ is selected: the UMAP embedding (Figure 1)
248 reveals three distinct manifold branches, and the entropy gain from $M=2$ to $M=3$ is substantial
249 (+0.30) while the gain from $M=3$ to $M=4$ is marginal (+0.18) at higher BIC cost.

Table 3: Regime count selection on 1D Burgers’. BIC is the EM reconstruction objective; H : mean assignment entropy; $\bar{\varepsilon}$: mean relative L^2 test error over $\nu \in \{0.001, 0.1, 1.0\}$.

M	BIC ($\times 10^4$)	H	$\bar{\varepsilon}$
2	3.90	0.622	0.209
3	4.92	0.920	0.222
4	5.47	1.099	0.228




Figure 1: TEMPO regime structure on 1D Burgers’ ($M=3$, 19 EM iterations). *Left*: UMAP coloured by regime assignment. *Centre*: same embedding coloured by ν ; each branch spans multiple viscosity values. *Right*: GMM fit in POD coefficient space (α_1, α_2) ; stars mark regime centroids μ_m .

Gating structure. Table 4 reports the mean EM assignment weight $\bar{\gamma}_m^\nu = |\mathcal{T}_\nu|^{-1} \sum_{i \in \mathcal{T}_\nu} \gamma_{im}^*$ for each viscosity class, alongside the global mixture weights π_m . The per- ν weights are nearly identical to π_m across all three regimes, meaning each viscosity class is split between regimes in almost exactly the same proportion as the full dataset. This confirms that regime membership is determined by the shape of the initial condition u_0 , not by the viscosity label ν directly, consistent with the UMAP structure in Figure 1.

Table 4: Mean EM assignment weights per viscosity (1D Burgers’, TEMPO(POD), $M=3$) compared to global mixture weights π_m . Each data row sums to 1. The per- ν weights closely match π_m , confirming that regime membership is driven by the shape of u_0 , not by ν .

ν	Regime 1	Regime 2	Regime 3
0.001	0.32	0.28	0.40
0.1	0.32	0.26	0.42
1.0	0.33	0.26	0.41
π_m	0.32	0.27	0.41

Regime structure. Figure 1 shows the UMAP embedding of all training snapshots, coloured by regime assignment and by ν . The three discovered regimes do *not* coincide with individual viscosity values: each manifold branch spans multiple ν , confirming that the EM partition reflects intrinsic solution geometry rather than the physical parameter label. After convergence, mixture weights are $\pi = (0.324, 0.267, 0.409)$, yielding a balanced, non-degenerate partition.

4.4 Reconstruction Quality

Figure 2 shows a representative space–time prediction from TEMPO(NeuralPOD) on 1D Burgers’ ($\nu=1.0$). The prediction closely tracks the ground truth with small, uniformly distributed residuals and no systematic error near the diffusive front. Additional reconstructions and error distributions for all baselines are provided in Appendix A.2.

5 Conclusion

Operator learning typically uses a single global representation across the full parameter space. When the governing parameter drives qualitative change in solution structure, a single basis must fit incompatible geometries simultaneously, concentrating error where accurate prediction matters most.

TEMPO addresses this by treating regime membership as a latent variable. EM discovers M regimes from unlabelled snapshots and fits a dedicated basis to each, without regime supervision. On 1D Burgers’ across three orders of viscosity, the discovered regimes cross viscosity boundaries, confirming that the partition reflects solution geometry rather than the physical parameter directly. Regime-specific bases reduce reconstruction error relative to a global basis at equal mode count. The gated online operator outperforms a single joint baseline across the full parameter range.




Figure 2: Space-time reconstruction on 1D Burgers' ($\nu=1.0$, TEMPO(NeuralPOD)): ground truth (left), prediction (centre), and point-wise absolute error (right).

Several directions remain open. Conditioning the basis functions on κ would allow each mode to encode parameter-dependent structure directly. Extending TEMPO to higher-dimensional benchmarks will test whether regime discovery generalises beyond one-dimensional shock dynamics. Convergence guarantees for Generalized EM under nonlinear basis parameterisation remain an open theoretical problem.

References

- [1] Nikola B. Kovachki, Samuel Lanthaler, and Andrew M. Stuart. Operator learning: Algorithms and analysis. *arXiv*, 2024.
- [2] Lu Lu, Pengzhan Jin, and George Em Karniadakis. Deeponet: Learning nonlinear operators for identifying differential equations based on the universal approximation theorem of operators. *arXiv*, 2019.
- [3] Lu Lu, Xuhui Meng, Shengze Cai, Zhiping Mao, Somdatta Goswami, Zhongqiang Zhang, and George Em Karniadakis. A comprehensive and fair comparison of two neural operators (with practical extensions) based on fair data. *arXiv*, 2022.
- [4] Ramansh Sharma and Varun Shankar. Ensemble and mixture-of-experts deeponets for operator learning. *arXiv*, 2025.
- [5] Changhong Mou, Binghang Lu, and Guang Lin. Neural-pod: A plug-and-play neural operator framework for infinite-dimensional functional nonlinear proper orthogonal decomposition. *arXiv*, 2026.
- [6] David Amsallem, Matthew J. Zahr, and Charbel Farhat. Nonlinear model order reduction based on local reduced-order bases. *International Journal for Numerical Methods in Engineering*, 92(10):891–916, 2012. doi: 10.1002/nme.4371.
- [7] Martin Hess, Alessandro Alla, Annalisa Quaini, Gianluigi Rozza, and Max Gunzburger. A localized reduced-order modeling approach for PDEs with bifurcating solutions. *Computer Methods in Applied Mechanics and Engineering*, 351:379–403, 2019. doi: 10.1016/j.cma.2019.03.050.
- [8] Thomas Daniel, Fabien Casenave, Nissrine Akkari, and David Ryckelynck. Model order reduction assisted by deep neural networks (ROM-net). *Advanced Modeling and Simulation in Engineering Sciences*, 7(1):16, 2020. doi: 10.1186/s40323-020-00153-6.
- [9] Matthew Tancik, Pratul P. Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan T. Barron, and Ren Ng. Fourier features let networks learn high frequency functions in low dimensional domains. In *Advances in Neural Information Processing Systems*, volume 33, pages 7537–7547, 2020.
- [10] Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred A. Hamprecht, Yoshua Bengio, and Aaron Courville. On the spectral bias of neural networks. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 5301–5310, 2019.

- 313 [11] Radford M. Neal and Geoffrey E. Hinton. A view of the EM algorithm that justifies in-
314 cremental, sparse, and other variants. In Michael I. Jordan, editor, *Learning in Graphical*
315 *Models*, volume 89 of *NATO ASI Series*, pages 355–368. Springer, Dordrecht, 1998. doi:
316 10.1007/978-94-011-5014-9_12.
- 317 [12] Makoto Takamoto, Timothy Praditia, Raphael Leiteritz, Dan MacKinlay, Francesco Alesiani,
318 Dirk Pflüger, and Mathias Niepert. Pdebench: An extensive benchmark for scientific machine
319 learning. *arXiv*, 2022.

Algorithm 1 WEIGHTEDNEURALPOD($\{s^{(i)}, \gamma_{im}\}_{i=1}^N, \tau$)

- 1: $\bar{s}_m \leftarrow \sum_{i=1}^N \gamma_{im} s^{(i)}$ ▷ responsibility-weighted mean trajectory
 - 2: Train $\phi_0^{(m)}$ by minimizing $\sum_j w_j (\phi_0^{(m)}(y_j) - (\bar{s}_m)_j)^2$
 - 3: $r^{(0,i)} \leftarrow s^{(i)} - \phi_0^{(m)}(y)$ for all i ; $k \leftarrow 0$
 - 4: $V_m \leftarrow \sum_{i=1}^N \gamma_{im} \|s^{(i)} - \bar{s}_m\|_w^2$ ▷ initial unmodeled variance
 - 5: **while** $\sum_{i=1}^N \gamma_{im} \|r^{(k,i)}\|_w^2 \geq \tau V_m$ **and** $k < K_{\max}$ **do**
 - 6: Solve γ_{im} -weighted (10) for $\phi_{k+1}^{(m)}$ and $\{\lambda_{k+1}^{(m,i)}\}$
 - 7: $r^{(k+1,i)} \leftarrow r^{(k,i)} - \lambda_{k+1}^{(m,i)} \phi_{k+1}^{(m)}(y)$ for all i
 - 8: $k \leftarrow k + 1$
 - 9: **end while**
 - 10: **return** $K_m \leftarrow k$, $\phi_0^{(m)}$, $\{\phi_k^{(m)}\}$, $\{\lambda_k^{(m,i)}\}_{i=1}^N\}_{k=1}^{K_m}$
-

Algorithm 2 TEMPO — Offline Phase

Require: $\mathcal{D} = \{u_0^{(i)}, \kappa^{(i)}, s^{(i)}\}_{i=1}^N$; $M, \varepsilon_s, \varepsilon_l, \varepsilon_p, \varepsilon_c, \tau$
Ensure: Regime bases $\{\Phi_m^*\}$, amplitudes $\{\Lambda_m^*\}$, responsibilities $\{\gamma_{im}^*\}$, weights $\{\pi_m^*\}$

Initialization

- 1: Compute global POD on $\{s^{(i)}\}$; project to $\alpha^{(i)} \in \mathbb{R}^P$ for all i
- 2: Run GMM-EM on $\{\alpha^{(i)}\} \rightarrow \gamma_{im}^{(0)}, \pi_m^{(0)}, \mu_m^{(0)}, \Sigma_m^{(0)}$
- 3: **for** $m = 1, \dots, M$ **do**
- 4: $\Phi_m^{(0)}, \Lambda_m^{(0)} \leftarrow \text{WEIGHTEDNEURALPOD}(\{s^{(i)}, \gamma_{im}^{(0)}\}, \tau)$ ▷ or weighted SVD for TEMPO(POD), see §3.1
- 5: **end for**
- 6: $\sigma^2 \leftarrow (2N)^{-1} \sum_{i,m} \gamma_{im}^{(0)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2$ (14)

EM iterations

- 7: **repeat**
- 8: *E-step*
- 9: **for** all $i = 1, \dots, N$ and $m = 1, \dots, M$ **do**
- 10: Compute $\hat{s}_m^{(i)}$ via (4) using $\Phi_m^{(q-1)}$ and $\Lambda_m^{(q-1)}$
- 11: $\gamma_{im}^{(q)} \leftarrow \text{Eq. (15)}$
- 12: **end for**
- 13: *M-step: mixture weights*
- 14: $\pi_m^{(q)} \leftarrow N_m^{(q)} / N$ for all m (16)
- 15: *M-step: bases*
- 16: **for** $m = 1, \dots, M$ **do**
- 17: Compute $\mu_m^{(q)}$ (18), $\Sigma_m^{(q)}$ (19), $\Delta_m^{(q)}$ (20)
- 18: **if** $\Delta_m^{(q)} < \varepsilon_s$ **then**
- 19: $\Phi_m^{(q)} \leftarrow \Phi_m^{(q-1)}$ ▷ Skip
- 20: **else if** $\Delta_m^{(q)} < \varepsilon_l$ **then**
- 21: Recompute $r_m^{(K_m, i)}$ and $V_m^{(q)} \leftarrow \sum_i \gamma_{im}^{(q)} \|s^{(i)} - \bar{s}_m^{(q)}\|_w^2$ ▷ Incremental
- 22: **if** $\sum_i \gamma_{im}^{(q)} \|r_m^{(K_m, i)}\|_w^2 \geq \tau \cdot V_m^{(q)}$ **then**
- 23: Train mode $K_m + 1$; $K_m \leftarrow K_m + 1$
- 24: **end if**
- 25: **if** $\|\phi_{K_m}^{(m)}\|_w^2 \cdot \sum_{i=1}^N \gamma_{im}^{(q)} (\lambda_{K_m}^{(m, i)})^2 < \varepsilon_p$ **then**
- 26: $K_m \leftarrow K_m - 1$
- 27: **end if**
- 28: **else**
- 29: $\Phi_m^{(q)}, \Lambda_m^{(q)} \leftarrow \text{WEIGHTEDNEURALPOD}(\{s^{(i)}, \gamma_{im}^{(q)}\}, \tau)$ ▷ Full rerun
- 30: **end if**
- 31: **end for**
- 32: **until** $\max_m \Delta_m^{(q)} < \varepsilon_c$
- 33: **return** $\Phi_m^* \leftarrow \Phi_m^{(q)}, \Lambda_m^* \leftarrow \Lambda_m^{(q)}, \gamma_{im}^* \leftarrow \gamma_{im}^{(q)}, \pi_m^* \leftarrow \pi_m^{(q)}$

322 A.2 Additional Experiments

323 **POD basis quality (1D Burgers’).** Figure 3 shows the singular value spectrum and cumulative
 324 explained variance for the global POD on 1D Burgers’. $K=30$ modes capture 99.99% of the total
 325 variance, achieving a Phase 1 relative L^2 reconstruction error of 1.47%.

326 **NeuralPOD basis quality (1D Burgers’).** Figure 4 shows the greedy residual norm as a function
 327 of mode count; the steep initial decay confirms that the Fourier NeuralPOD basis concentrates vari-
 328 ance efficiently. Figure 5 shows the learned spatial modes alongside representative reconstructions;
 329 modes are smooth and globally supported for large ν , and sharply localised near the shock front for
 330 small ν .




Figure 3: Global POD basis on 1D Burgers' ($\nu=1.0$, $K=30$ modes). *Left*: singular value spectrum. *Right*: cumulative explained variance; $K=30$ modes attain 99.99%.




Figure 4: Fourier NeuralPOD basis on 1D Burgers' ($\nu=1.0$): weighted residual norm versus greedy mode count.

331 **Baseline training curves.** Figures 6 and 7 show the Phase 2 MSE training curves for POD-
 332 DeepONet and NeuralPOD-DeepONet on 1D Burgers' ($\nu=1.0$).

333 **Error distributions.** Figures 8 and 9 show the full test-set error distributions for POD-DeepONet
 334 and NeuralPOD-DeepONet on 1D Burgers' ($\nu=1.0$).

335 **Representative reconstructions.** Figures 10 and 11 show representative space-time predictions
 336 for POD-DeepONet and NeuralPOD-DeepONet on 1D Burgers' ($\nu=1.0$).




Figure 5: Learned Fourier NeuralPOD basis functions and representative reconstructions on 1D Burgers' ($\nu=1.0$).




Figure 6: Phase 2 MSE training curve for POD-DeepONet on 1D Burgers' ($\nu=1.0$).



Figure 7: Phase 2 MSE training curve for NeuralPOD-DeepONet on 1D Burgers' ($\nu=1.0$).




Figure 8: POD-DeepONet test-set error distribution on 1D Burgers' ($\nu=1.0$). *Left*: histogram with mean (red dashed) and median (blue dashed). *Right*: sorted relative L^2 errors.




Figure 9: NeuralPOD-DeepONet test-set error distribution on 1D Burgers' ($\nu=1.0$). Layout as in Figure 8.




Figure 10: POD-DeepONet: representative reconstructions on 1D Burgers' ($\nu=1.0$). Each row: ground truth (*left*), prediction (*centre*), point-wise absolute error (*right*).




Figure 11: NeuralPOD-DeepONet: representative reconstructions on 1D Burgers' ($\nu=1.0$). Layout as in Figure 10.