
TEMPO: Regime-Aware Operator Learning with Locally Adapted POD Bases

Anonymous Author(s)

Affiliation

Address

email

Abstract

Operator learning methods typically approximate the solution map of a physical system with a single global model—an assumption that fails when solutions change character fundamentally across the parameter space.

We propose TEMPO: a framework that identifies latent dynamical regimes in snapshot data via Expectation-Maximization, builds a dedicated basis for each, and learns to route new inputs to the right regime at inference time. Models trained on a single parameter value fail badly on others—errors grow by orders of magnitude out-of-distribution. A single jointly-trained model partially addresses this but ignores underlying structure. TEMPO instead learns regime-specific bases jointly across all parameter values, outperforming the global joint baseline while recovering structure driven by both solution geometry and the physical parameter.

1 Introduction

Operator learning approximates the solution map of a parametric PDE from data, enabling instant prediction without re-solving the governing equations at each new parameter value [1]. The dominant paradigm—exemplified by DeepONet [2] and POD-DeepONet [3]—trains a single global model over the entire parameter space. This works well when solutions form a simple manifold, but breaks down when the governing parameter drives *qualitatively distinct* dynamics: viscosity separating laminar diffusion from near-discontinuous shocks, or source magnitude producing solutions that differ by orders of magnitude in amplitude. In these multi-regime settings, a single global basis must compromise between incompatible geometries, concentrating error precisely where accurate prediction matters most.

Several methods improve expressivity while retaining a single global representation. POD-PoU [4] trains a partition-of-unity mixture of POD bases with a shared branch, reducing error on sharp-gradient problems but without discovering the underlying regime structure. NeuralPOD [5] constructs nonlinear orthogonal basis functions via greedy residual minimization, representing each mode as a neural network and thus overcoming the linear subspace constraint. Both, however, operate on the full dataset globally and provide no operator for routing new inputs to the appropriate representation at inference.

Clustering snapshots and fitting local bases to each cluster is a classical idea in reduced-order modeling [6, 7]. Daniel et al. [8] extend this to deep learning by clustering solution subspaces on the Grassmann manifold and training a dictionary selector at inference. These methods use hard cluster assignments, operate offline on fixed datasets, and do not couple regime discovery with the operator learning objective—leaving basis quality and assignment accuracy disconnected.

We propose **TEMPO** (*Trajectory EM Mixture of POD Operators*), which closes this gap in a principled probabilistic framework. TEMPO models the trajectory ensemble as a Gaussian mixture and

recovers regime structure via Expectation-Maximization (EM): the E-step softly assigns each trajectory to regimes based on reconstruction quality, while the M-step rebuilds a dedicated NeuralPOD basis per regime so as bases sharpen, assignments improve, and vice versa. After convergence, a GatingNet and per-regime BranchNets are trained jointly, routing new inputs to the correct regime at inference.

Contributions:

- A Gaussian mixture model with EM that discovers M latent regimes and fits a dedicated NeuralPOD basis per regime, requiring no regime labels.
- NeuralPOD-DeepONet: the first deployable operator built on the NeuralPOD basis of [5], pairing it with a branch network that predicts mode coefficients for unseen inputs.
- An adaptive basis update rule based on a distribution-shift criterion, which determines when each regime’s basis warrants retraining and avoids redundant computation across EM iterations.
- On 1D Burgers’ and 2D Darcy flow, per-parameter specialists achieve low in-distribution error but fail catastrophically outside their training value; TEMPO(POD) trained jointly achieves lower average error across the full parameter range than any single-parameter specialist and outperforms a jointly-trained model without regime discovery, with discovered regimes reflecting both solution geometry and physical parameter structure.

2 Problem Statement

Let $\Omega \subset \mathbb{R}^d$ be a bounded spatial domain and $\mathcal{T} \subset \mathbb{R}_{\geq 0}$ a time interval.

We are given a dataset of N solution trajectories:

$$\mathcal{D} = \{ (u_0^{(i)}, \kappa^{(i)}, s^{(i)}) \}_{i=1}^N, \quad s^{(i)} = s(\cdot, \cdot; u_0^{(i)}, \kappa^{(i)}) \in L^2(\Omega \times \mathcal{T}), \quad (1)$$

where $u_0^{(i)} \in \mathcal{U} = L^2(\Omega)$ is the initial condition and $\kappa^{(i)} \in \mathcal{K}$ the physical parameter. Each trajectory $s^{(i)}$ is discretized on $N_y = N_t N_x$ quadrature points $\{y_j = (x_j, t_j), w_j\}_{j=1}^{N_y}$ covering the product grid $\Omega \times \mathcal{T}$, with weighted norm $\|f\|_w^2 = \sum_j w_j f(y_j)^2$.

Definition 1 (Solution operator). *The solution operator*

$$\mathcal{G} : \mathcal{U} \times \mathcal{K} \rightarrow L^2(\Omega \times \mathcal{T})$$

maps each pair (u_0, κ) to the full space-time solution: $\mathcal{G}(u_0, \kappa) = s(\cdot, \cdot; u_0, \kappa)$.

The operator $\mathcal{G}$ is unknown and expensive to evaluate directly. The *operator learning problem* is to construct $\hat{\mathcal{G}} \approx \mathcal{G}$ from $\mathcal{D}$ alone, generalizing to unseen $(u_0, \kappa) \notin \mathcal{D}$ without solving the PDE.




Figure 1: Training dataset $\mathcal{D}$: solution trajectories $s^{(i)}$ on $\Omega \times \mathcal{T}$ for different $\kappa^{(i)} \in \mathcal{K}$. Solutions change character fundamentally across parameter values.

Multi-regime assumption. When solutions change character fundamentally across $\mathcal{K}$, the optimal low-dimensional representation differs from regime to regime, and no single global basis captures the full solution diversity. We formalize this observation via a latent variable model over the trajectory ensemble.

Definition 2 (Generative model). *Each trajectory arises from one of M latent regimes. The regime label $z_i \in \{1, \dots, M\}$ is unobserved and drawn from a categorical prior,*

$$z_i \sim \text{Categorical}(\pi_1, \dots, \pi_M), \quad \pi_m \geq 0, \quad \sum_m \pi_m = 1. \quad (2)$$

Conditioned on $z_i = m$, the trajectory is distributed as

$$(s^{(i)} \mid z_i = m) \sim \mathcal{N}(\hat{s}_m^{(i)}, \sigma^2 W^{-1}), \quad (3)$$

where $W = \text{diag}(w_1, \dots, w_{N_y})$ is the quadrature weight matrix, $\sigma^2 > 0$ controls the softness of regime assignments, and $\hat{s}_m^{(i)}$ is the regime- m approximation to $s^{(i)}$.

The generative model leaves $\hat{s}_m^{(i)}$ unspecified. We require only that it admits a *low-rank trajectory decomposition*:

$$\hat{s}_m^{(i)}(y) = \phi_0^{(m)}(y) + \sum_{k=1}^{K_m} \lambda_k^{(m,i)} \phi_k^{(m)}(y), \quad y = (x, t) \in \Omega \times \mathcal{T}, \quad (4)$$

where $\phi_k^{(m)} : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$ are spatiotemporal mode functions and $\lambda_k^{(m,i)} \in \mathbb{R}$ a scalar amplitude for trajectory i on mode k . This form encompasses POD [3] and learned neural bases [5]; in this work we adopt the latter.

Definition 3 (Regime approximator). *The parameters of regime m are $\Phi_m = \{\phi_k^{(m)}\}_{k=0}^{K_m}$ and scalar amplitudes $\Lambda_m = \{\lambda_k^{(m,i)}\}_{k=1, i=1}^{K_m, N}$, where $\phi_0^{(m)} : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$ is the regime mean field, $\{\phi_k^{(m)}\}_{k \geq 1}$ spatiotemporal mode functions, $\lambda_k^{(m,i)} \in \mathbb{R}$ the scalar amplitude of trajectory i on mode k , and $K_m \leq K_{\max}$ the adaptive mode count.*

The learning problem decomposes into two phases: discovering the latent regime structure from $\mathcal{D}$ *offline*, then training a deployable operator for fast inference *online*.

Phase 1: Offline regime discovery. Since the regime labels z_i in (3) are unobserved, we recover the bases $\{\Phi_m\}$, mixture weights π , and noise level σ^2 by maximizing the observed-data log-likelihood:

$$(\Phi^*, \pi^*, \sigma^{2*}) = \arg \max_{\Phi, \pi, \sigma^2} \sum_{i=1}^N \log \sum_{m=1}^M \frac{\pi_m \det(W)^{1/2}}{(2\pi\sigma^2)^{N_y/2}} \exp\left(-\frac{\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2}{2\sigma^2}\right). \quad (5)$$

The solution yields soft regime responsibilities $\gamma_{im}^* = p(z_i = m \mid s^{(i)}, \Phi^*, \pi^*, \sigma^{2*})$ that serve as supervision for Phase 2.

Phase 2: Online operator learning. With bases $\{\Phi_m^*\}$ and responsibilities $\{\gamma_{im}^*\}$ frozen, we train a gating network (*GatingNet*) and M branch networks (*BranchNets*) that map a new input (u_0, κ) to a weighted combination of local regime predictions. The precise architecture and loss are given in Section 3.4. The quality of the final operator is measured by the mean relative L^2 error:

$$\varepsilon = \frac{1}{N_{\text{test}}} \sum_{i=1}^{N_{\text{test}}} \frac{\|s^{(i)} - \hat{s}(\cdot; u_0^{(i)}, \kappa^{(i)})\|_w}{\|s^{(i)}\|_w}. \quad (6)$$

3 Method

All methods studied in this work share a common two-phase structure: an *offline* basis-fitting phase that constructs a low-rank approximation to the solution manifold from snapshots in $\mathcal{D}$, and an *online* operator-learning phase that trains a branch network to predict expansion coefficients from a new input (u_0, κ) . The four methods differ along two independent axes:




Figure 2: **TEMPO pipeline.** Phase 1 (offline): snapshots are soft-assigned to M regimes via EM, and a dedicated basis $\{\Phi_m\}$ is fitted per regime. Phase 2 (online): a GatingNet routes each input to regime weights; M BranchNets predict expansion coefficients; the output is a weighted sum of local reconstructions.

- **Basis type.** A *linear* (POD) basis uses SVD of the (weighted) snapshot matrix; a *nonlinear* (NeuralPOD) basis learns spatiotemporal mode networks by sequential residual minimization.
- **Number of regimes.** A *single-regime* operator fits one global basis to all snapshots; a *multi-regime* operator (TEMPO) discovers M latent regimes via EM and assigns a dedicated basis to each.

This gives four combinations: POD-DeepONet [3] (linear, global), NeuralPOD-DeepONet (nonlinear, global, our extension of [5]), TEMPO(POD) (linear, M regimes), and TEMPO(NeuralPOD) (nonlinear, M regimes).

3.1 Basis Representations

3.1.1 POD Basis

Let $\phi_0 := \bar{s} = \frac{1}{N} \sum_i s^{(i)}$ be the mean field and $\tilde{S} \in \mathbb{R}^{N_y \times N}$ the matrix with columns $\tilde{s}^{(i)} = s^{(i)} - \phi_0$. The rank- K truncated SVD $\tilde{S} \approx U_K \Sigma_K V_K^\top$ yields orthonormal modes $\{\phi_k\}_{k=1}^K$ (columns

of U_K) and coefficients $\lambda_k^{(i)} = [\Sigma_K V_K^\top]_{ki}$; the approximation

$$\hat{s}^{(i)} = \phi_0 + \sum_{k=1}^K \lambda_k^{(i)} \phi_k \quad (7)$$

is optimal in the L^2 sense among all rank- K affine representations.

For TEMPO(POD), the M-step (17) is solved analytically: set $\phi_0^{(m)} := \bar{s}_m = \frac{1}{N_m} \sum_i \gamma_{im} s^{(i)}$, where $N_m = \sum_i \gamma_{im}$, and apply SVD to the weighted centered matrix with columns $\sqrt{\gamma_{im}} (s^{(i)} - \phi_0^{(m)})$. This is the closed-form linear analogue of Algorithm 1.

3.1.2 NeuralPOD Basis

To overcome the linearity constraint, we parametrize each mode function $\phi_k : \Omega \times \mathcal{T} \rightarrow \mathbb{R}$ as a Fourier feature network [9]; we call the resulting basis a *Fourier NeuralPOD* basis:

$$\phi_k(y) = f_{\theta_k}(\psi(y)), \quad \psi(y) = [\sin(2\pi B y), \cos(2\pi B y)], \quad (8)$$

where $B \in \mathbb{R}^{F \times d}$ is a fixed random matrix drawn once at initialization and $f_{\theta_k} : \mathbb{R}^{2F} \rightarrow \mathbb{R}$ is a small MLP with Tanh activations. Plain MLPs exhibit spectral bias [10]; the Fourier encoding resolves this. For multi-scale coverage, the rows of B are partitioned evenly across L frequency scales $\sigma_1 < \dots < \sigma_L$: the ℓ -th block is drawn from $\mathcal{N}(0, \sigma_\ell^2 I)$. The mean field ϕ_0 uses the same architecture.

The scalar amplitudes $\lambda_k^{(i)} \in \mathbb{R}$ are direct learnable parameters, one per training trajectory, optimized jointly with ϕ_k . After Phase 1 they are frozen and serve as regression targets for the BranchNet in Phase 2.

Modes are learned by *greedy residual minimization* [5]. Let $r^{(0,i)} = s^{(i)} - \phi_0$, where ϕ_0 is fitted to the snapshot mean $\bar{s}$ defined above. For each subsequent mode $k \geq 1$, given the residual

$$r^{(k-1,i)}(y) = s^{(i)}(y) - \phi_0(y) - \sum_{\ell=1}^{k-1} \lambda_\ell^{(i)} \phi_\ell(y), \quad (9)$$

we solve

$$\min_{\phi_k, \{\lambda_k^{(i)}\}} \sum_{i=1}^N \sum_j w_j \left(\lambda_k^{(i)} \phi_k(y_j) - r_j^{(k-1,i)} \right)^2 \quad (10)$$

(the global case with uniform weights; the γ -weighted generalisation used in TEMPO is Algorithm 1). Mode addition stops when the residual falls below fraction τ of the initial variance,

$$\sum_{i=1}^N \|r^{(k,i)}\|_w^2 < \tau \cdot \sum_{i=1}^N \|s^{(i)} - \bar{s}\|_w^2, \quad (11)$$

or $k = K_{\max}$. Each mode captures the largest remaining variance component the nonlinear analogue of Gram-Schmidt orthogonalization.

3.2 Single-Regime Operators

3.2.1 POD-DeepONet

POD-DeepONet [3] replaces the learned trunk network of DeepONet [2] with the fixed POD basis. After computing the global modes $\{\phi_k\}_{k=1}^K$ and training coefficients $\{\lambda^{(i)}\}$ offline, a branch network $\mathcal{B}_\theta : \mathbb{R}^{m_s + d_\kappa} \rightarrow \mathbb{R}^K$ is trained online by regression against these coefficients:

$$\mathcal{L}_{\text{branch}} = \frac{1}{N} \sum_{i=1}^N \|\mathcal{B}_\theta(u_0^{(i)}, \kappa^{(i)}) - \lambda^{(i)}\|^2. \quad (12)$$

The prediction at any $(x, t) \in \Omega \times \mathcal{T}$ is

$$\hat{s}(x, t; u_0, \kappa) = \phi_0(x, t) + \sum_{k=1}^K \mathcal{B}_{\theta,k}(u_0, \kappa) \phi_k(x, t). \quad (13)$$

3.2.2 NeuralPOD-DeepONet

Mou et al. [5] introduce NeuralPOD as a dimensionality reduction for PDE solutions, establishing the greedy basis-fitting procedure of Section 3.1.2. However, they do not construct a deployable solution operator: no network is proposed to predict the mode coefficients for an unseen initial condition. We close this gap by substituting the NeuralPOD basis into the POD-DeepONet framework; we call this combination *NeuralPOD-DeepONet*.

The procedure follows POD-DeepONet exactly with the Fourier NeuralPOD basis substituted for POD. Offline fitting yields coefficients $\lambda^{(i)} \in \mathbb{R}^K$; online, a branch network minimizes (12) and prediction follows (13). At equal mode count K , the nonlinear basis spans a strictly richer function space than POD.

3.3 TEMPO: Offline Regime Discovery

A single global basis whether linear or nonlinear must compromise between qualitatively distinct regimes, as the optimal basis geometry differs fundamentally between them. TEMPO avoids this by constructing a dedicated basis for each of M latent regimes, discovered from unlabelled trajectories via EM on the generative model of Definition 2.

Initialization. A global POD is computed on all N trajectories; each trajectory is projected to a P -dimensional coefficient vector $\alpha^{(i)} \in \mathbb{R}^P$. Running a standard GMM on $\{\alpha^{(i)}\}$ yields initial responsibilities $\gamma_{im}^{(0)}$, weights $\pi_m^{(0)}$, and per-regime statistics $(\mu_m^{(0)}, \Sigma_m^{(0)})$. Each regime basis $\Phi_m^{(0)}$ is then fitted to the responsibility-weighted ensemble: for TEMPO(NeuralPOD) via WEIGHTEDNEURALPOD (Algorithm 1); for TEMPO(POD) via the weighted SVD of Section 3.1. The noise level σ^2 is calibrated from the initial reconstructions rather than treated as a free hyperparameter:

$$\sigma^2 \leftarrow \frac{1}{2N} \sum_{i=1}^N \sum_{m=1}^M \gamma_{im}^{(0)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2, \quad (14)$$

which ties the softness of the E-step directly to the quality of the initial bases.

E-step. At iteration q , the responsibility of trajectory i for regime m is updated by Bayes' rule on the Gaussian likelihood (3):

$$\gamma_{im}^{(q)} = \frac{\pi_m^{(q-1)} \exp\left(-\frac{\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2}{2\sigma^2}\right)}{\sum_{j=1}^M \pi_j^{(q-1)} \exp\left(-\frac{\|s^{(i)} - \hat{s}_j^{(i)}\|_w^2}{2\sigma^2}\right)}, \quad (15)$$

where $\hat{s}_m^{(i)}$ is evaluated via (4) using $\Phi_m^{(q-1)}$ and $\Lambda_m^{(q-1)}$. As the bases improve in the M-step, the reconstructions $\hat{s}_m^{(i)}$ sharpen and the E-step automatically refines regime assignments in response distinguishing TEMPO from a GMM applied to a fixed feature space. When κ drives large amplitude variation (as with $\beta \in [0.01, 100]$ in the Darcy experiments), replacing the squared norm with the relative squared norm $\|s^{(i)} - \hat{s}_m^{(i)}\|_w^2 / \|s^{(i)}\|_w^2$ throughout — in (15), (14), and (17) — keeps E- and M-steps consistent under a heteroscedastic likelihood with noise variance proportional to $\|s^{(i)}\|_w^2$. When κ is observed, the GMM initialization can operate in $\log \kappa$ space, and an optional κ -prior term in the E-step anchors assignments to the known parameter structure.

M-step. With responsibilities fixed, the M-step gives a closed-form update for the mixture weights,

$$\pi_m^{(q)} \leftarrow \frac{N_m^{(q)}}{N}, \quad N_m^{(q)} = \sum_{i=1}^N \gamma_{im}^{(q)}, \quad (16)$$

and a weighted reconstruction problem for each basis:

$$\min_{\Phi_m, \Lambda_m} \sum_{i=1}^N \gamma_{im}^{(q)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2, \quad (17)$$

solved by weighted SVD for TEMPO(POD) or WEIGHTEDNEURALPOD (Algorithm 1) for TEMPO(NeuralPOD). Since gradient descent solves the NeuralPOD subproblem without guaranteeing global optimality, TEMPO is a Generalized EM algorithm [11], which guarantees that the observed-data log-likelihood is non-decreasing at every iteration.

Adaptive basis update. Retraining all M NeuralPOD bases at every EM iteration is computationally prohibitive. We introduce a distribution-shift criterion to determine whether each basis requires updating. Let $\alpha^{(i)}$ denote the global POD projection of $s^{(i)}$, computed once during initialization and fixed thereafter. The responsibility-weighted first and second moments of regime m are

$$\mu_m^{(q)} = \frac{1}{N_m^{(q)}} \sum_{i=1}^N \gamma_{im}^{(q)} \alpha^{(i)}, \quad (18)$$

$$\Sigma_m^{(q)} = \frac{1}{N_m^{(q)}} \sum_{i=1}^N \gamma_{im}^{(q)} (\alpha^{(i)} - \mu_m^{(q)})(\alpha^{(i)} - \mu_m^{(q)})^\top. \quad (19)$$

The distribution-shift criterion measures the relative change in these moments between consecutive iterations:

$$\Delta_m^{(q)} = \frac{\|\mu_m^{(q)} - \mu_m^{(q-1)}\|_2}{\|\mu_m^{(q-1)}\|_2 + \epsilon} + \frac{\|\Sigma_m^{(q)} - \Sigma_m^{(q-1)}\|_F}{\|\Sigma_m^{(q-1)}\|_F + \epsilon}. \quad (20)$$

where $\epsilon > 0$ is a small numerical stabiliser. $\Delta_m^{(q)}$ measures how much the composition of regime m has changed, not just the volume of reassignments. Two trajectories with very different $\alpha^{(i)}$ swapping regimes produce a large $\Delta_m^{(q)}$ and trigger a basis update; two similar snapshots doing so do not. Given thresholds $0 < \varepsilon_s < \varepsilon_l$, regime m is SKIPPED if $\Delta_m^{(q)} < \varepsilon_s$, updated INCREMENTALLY (add or prune one mode) if $\varepsilon_s \leq \Delta_m^{(q)} < \varepsilon_l$, and fully retrained from scratch (FULL RERUN) otherwise. In the INCREMENTAL case, we compute the residual of the current basis under the updated responsibilities $\gamma^{(q)}$, add a new mode if $\sum_i \gamma_{im}^{(q)} \|r_m^{(K_m, i)}\|_w^2 \geq \tau \cdot \sum_i \gamma_{im}^{(q)} \|s^{(i)} - \bar{s}_m^{(q)}\|_w^2$, and prune the last mode if its weighted contribution falls below ε_p . The FULL RERUN case uses a cold-start random initialization, as warm-starting may trap gradient descent in the geometry of the previous regime.

Convergence and model selection. EM terminates when $\max_m \Delta_m^{(q)} < \varepsilon_c$. The number of regimes $M \in \{2, 3, 4\}$ is selected by a two-criterion rule: BIC on the TEMPO EM log-likelihood (lower is better, penalises model complexity) and mean assignment entropy H (higher gain signals a genuinely new regime); see Table 4 (Appendix A.2). The complete offline procedure is given in Algorithm 2.

3.4 TEMPO: Online Gated Operator

After Algorithm 2 converges, all regime bases $\{\Phi_m^*\}$ and responsibilities $\{\gamma_{im}^*\}$ are frozen. The online phase trains a GatingNet $G : \mathbb{R}^{m_s+d_\kappa} \rightarrow \Delta^{M-1}$ that maps (u_0, κ) to regime weights, and M BranchNets $\mathcal{B}_m : \mathbb{R}^{m_s+d_\kappa} \rightarrow \mathbb{R}^{K_m}$ that predict expansion coefficients for each regime. The TEMPO prediction at query point $y = (x, t) \in \Omega \times \mathcal{T}$ is

$$\hat{s}(y; u_0, \kappa) = \sum_{m=1}^M g_m(u_0, \kappa) \left[\phi_0^{(m)}(y) + \sum_{k=1}^{K_m} b_{m,k}(u_0, \kappa) \phi_k^{(m)}(y) \right], \quad (21)$$

where $\mathbf{g} = G(u_0, \kappa)$ and $\mathbf{b}_m = \mathcal{B}_m(u_0, \kappa)$. The BranchNet outputs $\mathbf{b}_m$ predict the expansion coefficients $\lambda_m^{(i)}$ for new inputs (u_0, κ) not seen in training. The prediction is mesh-free and evaluable at any $y \in \Omega \times \mathcal{T}$ without retraining.

Training objective. Writing $g_m^{(i)} = g_m(u_0^{(i)}, \kappa^{(i)})$, the online loss is

$$\mathcal{L}_{\text{online}} = \underbrace{\frac{1}{N} \sum_{i=1}^N \|s^{(i)} - \hat{s}(\cdot; u_0^{(i)}, \kappa^{(i)})\|_w^2}_{\mathcal{L}_{\text{data}}} + \underbrace{\lambda_{\text{KL}} \frac{1}{N} \sum_{i,m} g_m^{(i)} \log \frac{g_m^{(i)}}{\gamma_{im}^*}}_{\mathcal{L}_{\text{KL}}} - \underbrace{\lambda_H \frac{1}{N} \sum_{i,m} g_m^{(i)} \log g_m^{(i)}}_{-\mathcal{L}_{\text{ent}}}. \quad (22)$$

$\mathcal{L}_{\text{data}}$ drives reconstruction accuracy. $\mathcal{L}_{\text{KL}}$ anchors the online gating to the offline regime partition: it penalizes deviation of $\mathbf{g}^{(i)}$ from the EM responsibilities γ_i^* , so the network cannot discard the discovered structure in pursuit of a simpler routing. $\mathcal{L}_{\text{ent}}$ prevents collapse to a single expert, ensuring all M regime bases contribute. The hyperparameters $\lambda_{\text{KL}}, \lambda_H \geq 0$ trade off these objectives.

4 Computational Experiments

We evaluate TEMPO on parametric PDE benchmarks, verifying three claims: (i) the offline EM phase recovers interpretable, physics-consistent regime structure from unlabelled snapshots; (ii) locally adapted bases achieve lower reconstruction error than a single global basis at equal mode count; (iii) a jointly-trained gated operator with regime-specific bases outperforms a single jointly-trained basis without regime discovery.

4.1 Experimental Setup

Datasets. We use three parametric PDE benchmarks; the first two are from PDEBench [12].

1D Burgers’ equation.

$$u_t + u u_x = \nu u_{xx}, \quad x \in (-1, 1), \quad t \in (0, 2], \quad (23)$$

with periodic boundary conditions and random initial conditions. Viscosity $\nu \in \{0.001, 0.01, 0.1, 1.0\}$ spans four qualitatively distinct solution regimes: smooth diffusive profiles at $\nu=1.0$ and near-discontinuous shock structures at $\nu=0.001$. Each value provides $N=9,500$ trajectories on $N_x=1,024$ spatial points and $N_t=201$ time steps (38,000 total).

2D Darcy flow.

$$-\nabla \cdot (a(x) \nabla u(x)) = \beta, \quad x \in (0, 1)^2, \quad (24)$$

with zero Dirichlet boundary conditions. Here $a(x)$ is a spatially varying permeability field sampled from a random sum of Gaussians, and $\beta \in \{0.01, 0.1, 1.0, 10.0, 100.0\}$ is the constant source term. Each value provides $N=8,000$ training and 1,000 test samples discretised on a 128×128 uniform grid ($N_y=16,384$ spatial points), 40,000 trajectories in total. The operator maps $a(x)$ to the pressure solution $u(x)$.

2D incompressible Navier-Stokes.

$$\partial_t \omega + \mathbf{u} \cdot \nabla \omega = \frac{1}{\text{Re}} \Delta \omega + f, \quad \mathbf{x} \in (0, 1)^2, \quad t \in (0, 2], \quad (25)$$

with periodic boundary conditions, $\omega = \nabla \times \mathbf{u}$, divergence-free $\mathbf{u}$ recovered via the stream function, and Kolmogorov-type forcing $f(\mathbf{x}) = 0.1 (\sin 2\pi(x_1+x_2) + \cos 2\pi(x_1+x_2))$ [13]. Reynolds number $\text{Re} \in \{100, 1000, 3600, 10000\}$ spans four qualitatively distinct regimes from smooth laminar flow to developed turbulence. Each value provides $N=5,000$ trajectories on a 64×64 periodic grid (20,000 total; 4,000/1,000 train/test); see Appendix A.1. Operator: $\mathbf{u}_0 \mapsto (\mathbf{u}_t)_{t=1}^{20}, \Delta t=0.1$.

Methods. We compare DeepONet [2], POD-DeepONet [3], FNO [13], and CNO [14] all trained jointly on all parameter values via a single global model. We introduce NeuralPOD-DeepONet (learned Fourier NeuralPOD basis with branch network); TEMPO(POD) and TEMPO(NeuralPOD) (regime-aware bases discovered via EM, routed by gating network). Per-parameter specialists serve as performance ceilings. All methods implemented by the authors. Evaluation: mean relative L^2 error (6) per parameter value.

4.2 Main Results

1D Burgers’ equation. Table 1 gives the cross-viscosity error matrix. Per- ν specialists are unviable across the full range: a $\nu=0.1$ model degrades from 0.048 to 0.461 at $\nu=0.001$. TEMPO(POD) addresses this by discovering regime structure jointly, outperforming the global joint baseline by 13–50% across all ν . TEMPO(NeuralPOD) recovers the same regimes but does not surpass the joint baseline, suggesting nonlinear basis fitting remains harder in the multi-regime setting.

2D Darcy flow. Table 2 gives the cross- β matrix. Specialist failure is catastrophic — Darcy solutions scale linearly with β , driving cross-parameter errors beyond 1000. TEMPO(POD) reduces error by 3–5 $\times$ over the joint baseline (e.g., $\beta=0.1$: 0.775 vs. 2.557; $\beta=100$: 0.073 vs. 0.347). TEMPO(NeuralPOD) recovers dramatically over its non-regime baseline ($\beta=0.1$: 3.813 vs. 72.9) but does not match TEMPO(POD). Recent methods (FNO, CNO) employ fundamentally different operator learning paradigms without explicit basis decomposition; TEMPO’s contribution is discovering and leveraging latent regime structure through interpretable, locally-adapted bases.

Table 1: **Relative L^2 test error on 1D Burgers’ equation.** Each per- ν specialist is trained on a single viscosity and evaluated on all three jointly-trained values; underlined entries are in-distribution. The large off-diagonal errors motivate joint training with regime discovery. All jointly-trained methods use $\nu \in \{0.001, 0.1, 1.0\}$. **Bold:** best per column among jointly-trained methods.

Method	Train ν	$\nu=0.001$	$\nu=0.1$	$\nu=1.0$
<i>Per-ν specialists (single-viscosity training)</i>				
DeepONet [2]	0.001	<u>0.536</u>	0.474	0.382
	0.1	0.535	<u>0.471</u>	0.349
	1.0	0.547	0.483	<u>0.283</u>
POD-DeepONet [3]	0.001	<u>0.273</u>	0.297	0.639
	0.01	0.274	0.274	0.614
	0.1	0.349	<u>0.048</u>	0.344
	1.0	0.461	0.260	<u>0.025</u>
NeuralPOD-DeepONet (<i>ours</i>)	0.001	<u>0.392</u>	0.288	0.535
	0.01	0.391	0.284	0.541
	0.1	0.426	<u>0.229</u>	0.371
	1.0	0.483	0.315	<u>0.183</u>
FNO [13]	0.001	<u>0.426</u>	0.529	0.623
	0.01	0.454	0.401	0.492
	0.1	0.650	<u>0.304</u>	0.521
	1.0	1.161	0.919	<u>0.158</u>
CNO [14] (2.0M)	0.01	0.687	0.955	2.344
	0.1	0.983	<u>0.199</u>	0.471
	1.0	1.870	1.084	<u>0.125</u>
<i>Joint training ($\nu \in \{0.001, 0.1, 1.0\}$)</i>				
DeepONet [2]	all	0.533	0.472	0.289
POD-DeepONet [3]	all	0.460	0.339	0.210
FNO [13]	all	0.427	0.298	0.202
CNO [14] (2.0M)	all	0.361	0.162	0.133
TEMPO(POD) (<i>ours</i>)	$M=3$	0.317	0.168	0.182
TEMPO(NeuralPOD) (<i>ours</i>)	$M=3$	0.542	0.456	0.288

4.3 Regime Discovery

We select $M=3$ based on entropy gain: H increases by +0.30 from $M=2$ to $M=3$ but only +0.18 from $M=3$ to $M=4$, indicating the third regime captures genuine structure that a fourth does not. BIC increases monotonically with M (favouring $M=2$ by complexity alone), so entropy is the decisive criterion here (Table 4, Appendix A.2). Figure 3 confirms three distinct manifold branches, one per discovered regime.

Each regime corresponds to a geometric type of initial condition u_0 that appears across all viscosity values. The per-viscosity gating weights (Table 5, Appendix A.2) are nearly uniform across $\nu \in \{0.001, 0.1, 1.0\}$, closely matching the global weights $\pi=(0.324, 0.267, 0.409)$: knowing ν alone does not predict regime membership. Within each regime, ν shapes the dynamics: the same initial geometry yields a shock at $\nu=0.001$ and a smooth wave at $\nu=1.0$. At inference, ν is passed to the branch network as a conditioning input.

Table 2: **Relative L^2 test error on 2D Darcy flow.** Each per- β specialist is trained on a single source value and evaluated on all five; underlined entries are in-distribution. Off-diagonal errors exceed 1 because the Darcy solution scales with β , so a specialist trained on $\beta=100$ predicts amplitudes $\sim 10^4$ times larger than the ground truth at $\beta=0.01$. TEMPO trains jointly on $\beta \in \{0.1, 1, 10, 100\}$; the $\beta=0.01$ regime is excluded from joint training (—). **Bold**: best per column among jointly-trained methods.

Method	Train β	$\beta=0.01$	$\beta=0.1$	$\beta=1.0$	$\beta=10$	$\beta=100$
<i>Per-β specialists (single-value training)</i>						
DeepONet [2]	0.01	<u>4.197</u>	0.615	0.944	0.994	0.999
	0.1	>10	<u>0.851</u>	0.875	0.987	0.999
	1.0	>10	>1	<u>0.433</u>	0.902	0.972
	10	>100	>10	>1	<u>0.243</u>	0.705
	100	>1000	>100	>10	>1	<u>0.078</u>
POD-DeepONet [3]	0.01	<u>0.449</u>	0.791	0.961	0.996	>1
	0.1	>1	<u>0.148</u>	0.871	0.986	0.999
	1.0	>10	>1	<u>0.048</u>	0.897	0.990
	10	>100	>10	>1	<u>0.035</u>	0.900
	100	>1000	>100	>10	>1	<u>0.042</u>
NeuralPOD-DeepONet (<i>ours</i>)	0.01	<u>1.407</u>	0.786	0.962	0.996	>1
	0.1	>1	<u>0.421</u>	0.902	0.989	0.999
	1.0	>10	>1	<u>0.135</u>	0.901	0.990
	10	>100	>10	>1	<u>0.050</u>	0.900
	100	>1000	>100	>10	>1	<u>0.037</u>
FNO [13]	0.01	<u>0.529</u>	0.791	0.967	0.997	0.9998
	0.1	6.26	<u>0.144</u>	0.877	0.988	0.999
	1.0	73.1	8.41	<u>0.034</u>	0.909	0.992
	10	>100	72.6	8.18	<u>0.013</u>	0.900
	100	>1000	256.3	43.2	6.13	<u>0.008</u>
CNO [14] (2.0M)	0.01	<u>0.388</u>	0.929	1.000	1.001	1.000
	0.1	10.6	<u>0.125</u>	0.948	1.004	1.002
	1.0	67.0	7.84	<u>0.024</u>	0.896	0.989
	10	>1000	99.1	8.37	<u>0.009</u>	0.898
<i>Joint training ($\beta \in \{0.1, 1, 10, 100\}$)</i>						
DeepONet [2]	all	—	76.53	8.503	1.352	0.493
POD-DeepONet [3]	all	—	2.557	0.363	0.181	0.347
NeuralPOD-DeepONet (<i>ours</i>)	all	—	72.9	6.92	0.357	0.345
FNO [13]	all	—	0.380	0.083	0.023	0.008
CNO [14] (2.0M)	all	—	0.133	0.026	0.008	0.005
TEMPO(POD) (<i>ours</i>)	$M=5$	—	0.775	0.178	0.136	0.073
TEMPO(NeuralPOD) (<i>ours</i>)	$M=5$	—	3.813	0.549	0.322	0.311




Figure 3: TEMPO(POD) regime structure on 1D Burgers' ($M=3$, 20 EM iterations). *Left*: UMAP coloured by regime. *Centre*: same embedding coloured by ν . *Right*: GMM fit in POD coefficient space (α_1, α_2); stars mark regime centroids μ_m .

Table 3: **Relative L^2 test error on 2D incompressible Navier-Stokes.** Each per-Re specialist is trained on a single Reynolds number and evaluated on all four; underlined entries are in-distribution. Joint methods train on all Re simultaneously. **Bold:** best per column among jointly-trained methods. (*pend.*): results pending.

Method	Train Re	Re=100	Re=1000	Re=3600	Re=10000	Infer. (ms)
<i>Per-Re specialists (single-Re training)</i>						
DeepONet [2]	100	<u>1.000</u>	1.002	1.003	1.004	—
	1000	<u>1.000</u>	<u>1.000</u>	1.000	1.000	
	3600	1.002	1.000	<u>1.000</u>	1.000	
	10000	1.001	1.000	1.000	<u>1.000</u>	
POD-DeepONet [3]	100	<u>0.082</u>	0.748	0.865	0.877	—
	1000	0.667	<u>0.129</u>	0.671	0.812	
	3600	0.660	<u>0.392</u>	<u>0.180</u>	0.623	
	10000	0.934	0.687	0.556	<u>0.345</u>	
NeuralPOD-DeepONet (<i>ours</i>)	1000	0.762	<u>0.198</u>	0.607	0.761	—
FNO [13]	100	<u>0.013</u>	0.262	0.368	0.437	1.50
	1000	<u>0.287</u>	<u>0.033</u>	0.138	0.227	
	3600	0.381	0.150	<u>0.039</u>	0.128	
	10000	0.461	0.260	0.126	<u>0.042</u>	
<i>Joint training (all Re)</i>						
POD-DeepONet [3]	all	0.092	0.132	0.141	0.145	—
FNO [13]	all	0.009	0.020	0.024	0.027	0.61
NeuralPOD-DeepONet (<i>ours</i>)	all			<i>pend.</i>		<i>pend.</i>
CNO [14]	all			<i>pend.</i>		<i>pend.</i>
TEMPO(POD) (<i>ours</i>)	$M=4$	0.105	0.153	0.167	0.181	0.05
TEMPO(NeuralPOD) (<i>ours</i>)	<i>pend.</i>			<i>pend.</i>		<i>pend.</i>

5 Conclusion

Operator learning typically uses a single global representation across the full parameter space. When the governing parameter drives qualitative change in solution structure, a single basis must fit incompatible geometries simultaneously, concentrating error where accurate prediction matters most.

TEMPO addresses this by treating regime membership as a latent variable. EM discovers M regimes from unlabelled snapshots and fits a dedicated basis to each, without regime supervision. Per-parameter specialists achieve low in-distribution error but are unusable across the full parameter range: errors grow by orders of magnitude out-of-distribution. TEMPO(POD) trained jointly outperforms a single joint baseline by 31–50% on Burgers’ and up to $3\times$ on Darcy flow, while discovered regimes span multiple parameter values, confirming that the partition reflects solution geometry and physical structure jointly. TEMPO(NeuralPOD) demonstrates that the regime discovery framework extends naturally to nonlinear bases; closing the gap with TEMPO(POD) in the joint training setting remains an open direction.

Several directions remain open. Conditioning the basis functions on κ would allow each mode to encode parameter-dependent structure directly. Extending TEMPO to higher-dimensional benchmarks will test whether regime discovery generalises beyond one-dimensional shock dynamics. Convergence guarantees for Generalized EM under nonlinear basis parameterisation remain an open theoretical problem.

References

- [1] Nikola B. Kovachki, Samuel Lanthaler, and Andrew M. Stuart. Operator learning: Algorithms and analysis. *arXiv*, 2024.
- [2] Lu Lu, Pengzhan Jin, and George Em Karniadakis. Deeponet: Learning nonlinear operators for identifying differential equations based on the universal approximation theorem of operators. *arXiv*, 2019.
- [3] Lu Lu, Xuhui Meng, Shengze Cai, Zhiping Mao, Somdatta Goswami, Zhongqiang Zhang, and George Em Karniadakis. A comprehensive and fair comparison of two neural operators (with practical extensions) based on fair data. *arXiv*, 2022.
- [4] Ramansh Sharma and Varun Shankar. Ensemble and mixture-of-experts deeponets for operator learning. *arXiv*, 2025.
- [5] Changhong Mou, Binghang Lu, and Guang Lin. Neural-pod: A plug-and-play neural operator framework for infinite-dimensional functional nonlinear proper orthogonal decomposition. *arXiv*, 2026.
- [6] David Amsallem, Matthew J. Zahr, and Charbel Farhat. Nonlinear model order reduction based on local reduced-order bases. *International Journal for Numerical Methods in Engineering*, 92(10):891–916, 2012. doi: 10.1002/nme.4371.
- [7] Martin Hess, Alessandro Alla, Annalisa Quaini, Gianluigi Rozza, and Max Gunzburger. A localized reduced-order modeling approach for PDEs with bifurcating solutions. *Computer Methods in Applied Mechanics and Engineering*, 351:379–403, 2019. doi: 10.1016/j.cma.2019.03.050.
- [8] Thomas Daniel, Fabien Casenave, Nissrine Akkari, and David Ryckelynck. Model order reduction assisted by deep neural networks (ROM-net). *Advanced Modeling and Simulation in Engineering Sciences*, 7(1):16, 2020. doi: 10.1186/s40323-020-00153-6.
- [9] Matthew Tancik, Pratul P. Srinivasan, Ben Mildenhall, Sara Fridovich-Keil, Nithin Raghavan, Utkarsh Singhal, Ravi Ramamoorthi, Jonathan T. Barron, and Ren Ng. Fourier features let networks learn high frequency functions in low dimensional domains. In *Advances in Neural Information Processing Systems*, volume 33, pages 7537–7547, 2020.
- [10] Nasim Rahaman, Aristide Baratin, Devansh Arpit, Felix Draxler, Min Lin, Fred A. Hamprecht, Yoshua Bengio, and Aaron Courville. On the spectral bias of neural networks. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 5301–5310, 2019.
- [11] Radford M. Neal and Geoffrey E. Hinton. A view of the EM algorithm that justifies incremental, sparse, and other variants. In Michael I. Jordan, editor, *Learning in Graphical Models*, volume 89 of *NATO ASI Series*, pages 355–368. Springer, Dordrecht, 1998. doi: 10.1007/978-94-011-5014-9_12.
- [12] Makoto Takamoto, Timothy Praditia, Raphael Leiteritz, Dan MacKinlay, Francesco Alesiani, Dirk Pflüger, and Mathias Niepert. Pdebench: An extensive benchmark for scientific machine learning. *arXiv*, 2022.
- [13] Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations. In *International Conference on Learning Representations*, 2021.
- [14] Bogdan Raonić, Roberto Molinaro, Jan Mishler, Tim Rohner, Francesca Bartolucci, Rima Alaifari, Siddhartha Mishra, and Emmanuel de Bézenac. Convolutional neural operators for robust and accurate learning of PDEs. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

A Appendix

A.1 2D Navier-Stokes Data Generation

We use the same equation and forcing as Li et al. [13] with a pseudo-spectral solver in stream-function formulation: diffusion treated implicitly (Crank–Nicolson), advection explicitly (second-order Adams–Bashforth), and $\frac{2}{3}$ -rule dealiasing on the 64×64 grid. Simulation time-step $\delta t = 0.01$; each recorded snapshot aggregates 10 sub-steps ($\Delta t = 0.1$). Initial vorticity fields are random sums of sinusoids,

$$\omega_0(\mathbf{x}) = \sum_{k=1}^7 \frac{a_k}{1+k} [\sin(2\pi k x_1 + \phi_k) + \cos(2\pi k x_2 + \phi_k)], \quad a_k \sim \mathcal{N}(0, 1), \quad \phi_k \sim \mathcal{U}[0, 2\pi].$$

A 5 s burn-in (500 simulation steps) discards transient behaviour before recording begins.

A.2 Additional Experiments

Regime model selection (1D Burgers’). Table 4 reports BIC, mean assignment entropy H , and mean test error $\bar{\varepsilon}$ for $M \in \{2, 3, 4\}$. Table 5 shows the per-viscosity EM responsibilities alongside global mixture weights π_m ; near-uniform agreement across ν confirms geometry-driven rather than parameter-driven regime membership.

Table 4: Regime count selection on 1D Burgers’. BIC = $-2 \log \mathcal{L} + k \log N$ (lower is better; penalises complexity); H : mean assignment entropy (higher gain signals a new genuine regime); $\bar{\varepsilon}$: mean relative L^2 test error. BIC increases monotonically, favouring $M=2$ by complexity alone. **Bold row**: selected configuration ($M=3$), chosen because the entropy gain $\Delta H = 0.30$ from $M=2 \rightarrow 3$ substantially exceeds $\Delta H = 0.18$ from $M=3 \rightarrow 4$.

M	BIC ($\times 10^4, \downarrow$)	H	$\bar{\varepsilon}$
2	3.90	0.622	0.209
3	4.92	0.920	0.222
4	5.47	1.099	0.228

Table 5: Mean EM assignment weights per viscosity (1D Burgers’, TEMPO(POD), $M=3$) vs. global mixture weights π_m . Each row sums to 1; near-uniform agreement across ν confirms regime membership is driven by u_0 geometry, not by ν .

ν	Regime 1	Regime 2	Regime 3
0.001	0.32	0.28	0.40
0.1	0.32	0.26	0.42
1.0	0.33	0.26	0.41
π_m	0.32	0.27	0.41

POD basis quality (1D Burgers’). Figure 4 shows the singular value spectrum and cumulative explained variance for the global POD on 1D Burgers’. $K=30$ modes capture 99.99% of the total variance, achieving a Phase 1 relative L^2 reconstruction error of 1.47%.

NeuralPOD basis quality (1D Burgers’). Figure 5 shows the greedy residual norm as a function of mode count; the steep initial decay confirms that the Fourier NeuralPOD basis concentrates variance efficiently. Figure 6 shows the learned spatial modes alongside representative reconstructions; modes are smooth and globally supported for large ν , and sharply localised near the shock front for small ν .

POD vs. NeuralPOD basis comparison (1D Burgers’, $\nu=1.0$). Figure 7 shows the singular-value decay of the global POD basis alongside the weighted residual decay of the Fourier NeuralPOD basis. Figure 8 shows the first K spatiotemporal mode functions for both bases; POD modes




Figure 4: Global POD basis on 1D Burgers' ($\nu=1.0$): singular value spectrum (*left*) and cumulative explained variance (*right*); $K=30$ modes attain 99.99%.




Figure 5: Fourier NeuralPOD basis on 1D Burgers' ($\nu=1.0$): weighted residual norm versus greedy mode count.




Figure 6: Learned Fourier NeuralPOD basis functions and representative reconstructions on 1D Burgers' ($\nu=1.0$).

are globally supported, while NeuralPOD modes adapt their spatial localisation to the data. Figure 9 shows the same comparison for $\nu=0.01$: NeuralPOD modes concentrate near the shock front while POD modes remain globally supported. Figure 10 compares per-sample reconstructions from both bases: ground truth, POD prediction, POD error, Fourier NeuralPOD prediction, and Fourier NeuralPOD error.




Figure 7: POD singular-value decay (*left*) and Fourier NeuralPOD residual decay (*right*) on 1D Burgers' ($\nu=1.0$).




Figure 8: First K basis modes on 1D Burgers' ($\nu=1.0$): POD (*top*) vs. Fourier NeuralPOD (*bottom*).




Figure 9: First K basis modes on 1D Burgers' ($\nu=0.01$): POD (*top*) vs. Fourier NeuralPOD (*bottom*).

Error distributions. Figures 11 and 12 show the full test-set error distributions for POD-DeepONet and NeuralPOD-DeepONet on 1D Burgers' ($\nu=1.0$).

Representative reconstructions. Figures 13 and 14 show representative space-time predictions for POD-DeepONet and NeuralPOD-DeepONet on 1D Burgers' ($\nu=1.0$).

TEMPO reconstructions (1D Burgers'). Figures 15–16 show representative space-time predictions from the two TEMPO variants on 1D Burgers'.

TEMPO regime structure (1D Burgers', NeuralPOD basis). Figure 17 shows the TEMPO(NeuralPOD) regime structure on 1D Burgers' ($M=3$). The partition mirrors




Figure 10: Phase 1 reconstruction comparison on 1D Burgers' ($\nu=1.0$). Each row shows one sample: ground truth, POD reconstruction, $|\text{POD error}|$, Fourier NeuralPOD reconstruction, $|\text{FNPOD error}|$.




Figure 11: POD-DeepONet test-set error distribution on 1D Burgers' ($\nu=1.0$): error histogram (*left*) and sorted errors (*right*).




Figure 12: NeuralPOD-DeepONet test-set error distribution on 1D Burgers' ($\nu=1.0$). Layout as in Figure 11.




Figure 13: POD-DeepONet: representative reconstructions on 1D Burgers' ($\nu=1.0$). Each row: ground truth (*left*), prediction (*centre*), point-wise absolute error (*right*).




Figure 14: NeuralPOD-DeepONet: representative reconstructions on 1D Burgers' ($\nu=1.0$). Layout as in Figure 13.




Figure 15: TEMPO(POD) on 1D Burgers' ($M=3$): ground truth (*left*), prediction (*centre*), point-wise absolute error (*right*).

TEMPO(POD) but with sharper assignments, reflecting the higher expressivity of the NeuralPOD basis.

TEMPO regime structure (2D Darcy flow, $M=4$). Figures 18–19 show the TEMPO Phase 1 regime structure on 2D Darcy flow ($M=4$) for POD and NeuralPOD bases. Regimes span multiple β values, confirming geometry-driven rather than parameter-driven partition.

TEMPO reconstructions (2D Darcy flow). Figures 20–21 show representative predictions from the two TEMPO variants on 2D Darcy flow.




Figure 16: TEMPO(NeuralPOD) on 1D Burgers' ($M=3$). Layout as in Figure 15.




Figure 17: TEMPO(NeuralPOD) on 1D Burgers' ($M=3$, $H=0.23$, cf. $H=0.84$ for POD). Layout as in Figure 3.




Figure 18: TEMPO(POD) on 2D Darcy flow ($M=4$, $H=0.75$, $\pi=(0.22, 0.28, 0.28, 0.22)$). Layout as in Figure 3.




Figure 19: TEMPO(NeuralPOD) on 2D Darcy flow ($M=4$, $H=0.68$). Layout as in Figure 3.




Figure 20: TEMPO(POD) on 2D Darcy flow ($M=5$): ground truth (*left*), prediction (*centre*), point-wise absolute error (*right*).




Figure 21: TEMPO(NeuralPOD) on 2D Darcy flow ($M=5$). Layout as in Figure 20.

A.3 Algorithms

Algorithm 1 WEIGHTEDNEURALPOD($\{s^{(i)}, \gamma_{im}\}_{i=1}^N, \tau$)

- 1: $\bar{s}^{(m)} \leftarrow \frac{1}{N_m} \sum_{i=1}^N \gamma_{im} s^{(i)}, \quad N_m = \sum_{i=1}^N \gamma_{im}$ $\triangleright$ responsibility-weighted mean trajectory
 - 2: Train $\phi_0^{(m)}$ by minimizing $\sum_j w_j (\phi_0^{(m)}(y_j) - \bar{s}_j^{(m)})^2$
 - 3: $r^{(0,i)} \leftarrow s^{(i)} - \phi_0^{(m)}(y)$ for all i ; $p \leftarrow 0$
 - 4: $V_m \leftarrow \sum_{i=1}^N \gamma_{im} \|s^{(i)} - \bar{s}^{(m)}\|_w^2$ $\triangleright$ initial unmodeled variance
 - 5: **while** $\sum_{i=1}^N \gamma_{im} \|r^{(p,i)}\|_w^2 \geq \tau V_m$ **and** $p < P_{\max}$ **do**
 - 6: Jointly minimize over $\phi_{p+1}^{(m)}$ and $\{\lambda_{p+1}^{(m,i)}\}_{i=1}^N$:

$$\sum_{i=1}^N \gamma_{im} \sum_j w_j \left(\lambda_{p+1}^{(m,i)} \phi_{p+1}^{(m)}(y_j) - r_j^{(p,i)} \right)^2$$
 - 7: $r^{(p+1,i)} \leftarrow r^{(p,i)} - \lambda_{p+1}^{(m,i)} \phi_{p+1}^{(m)}(y)$ for all i
 - 8: $p \leftarrow p + 1$
 - 9: **end while**
 - 10: **return** $P_m \leftarrow p, \phi_0^{(m)}, \{\phi_p^{(m)}, \{\lambda_p^{(m,i)}\}_{i=1}^N\}_{p=1}^{P_m}$
-

Algorithm 2 TEMPO — Offline Phase

Require: $\mathcal{D} = \{u_0^{(i)}, \kappa^{(i)}, s^{(i)}\}_{i=1}^N$; $M, \varepsilon_s, \varepsilon_l, \varepsilon_p, \varepsilon_c, \tau$
Ensure: Regime bases $\{\Phi_m^*\}$, amplitudes $\{\Lambda_m^*\}$, responsibilities $\{\gamma_{im}^*\}$, weights $\{\pi_m^*\}$

Initialization

- 1: Compute global POD on $\{s^{(i)}\}$; project to $\alpha^{(i)} \in \mathbb{R}^P$ for all i
- 2: Run GMM-EM on $\{\alpha^{(i)}\} \rightarrow \gamma_{im}^{(0)}, \pi_m^{(0)}, \mu_m^{(0)}, \Sigma_m^{(0)}$
- 3: **for** $m = 1, \dots, M$ **do**
- 4: $\Phi_m^{(0)}, \Lambda_m^{(0)} \leftarrow \text{WEIGHTEDNEURALPOD}(\{s^{(i)}, \gamma_{im}^{(0)}\}, \tau)$ ▷ or weighted SVD for TEMPO(POD), see §3.1
- 5: **end for**
- 6: $\sigma^2 \leftarrow (2N)^{-1} \sum_{i,m} \gamma_{im}^{(0)} \|s^{(i)} - \hat{s}_m^{(i)}\|_w^2$ (14)

EM iterations

- 7: **repeat**
- 8: *E-step*
- 9: **for** all $i = 1, \dots, N$ and $m = 1, \dots, M$ **do**
- 10: Compute $\hat{s}_m^{(i)}$ via (4) using $\Phi_m^{(q-1)}$ and $\Lambda_m^{(q-1)}$
- 11: $\gamma_{im}^{(q)} \leftarrow \text{Eq. (15)}$
- 12: **end for**
- 13: *M-step: mixture weights*
- 14: $\pi_m^{(q)} \leftarrow N_m^{(q)} / N$ for all m (16)
- 15: *M-step: bases*
- 16: **for** $m = 1, \dots, M$ **do**
- 17: Compute $\mu_m^{(q)}$ (18), $\Sigma_m^{(q)}$ (19), $\Delta_m^{(q)}$ (20)
- 18: **if** $\Delta_m^{(q)} < \varepsilon_s$ **then**
- 19: $\Phi_m^{(q)} \leftarrow \Phi_m^{(q-1)}$ SKIP
- 20: **else if** $\Delta_m^{(q)} < \varepsilon_l$ **then** INCREMENTAL
- 21: Compute $r_m^{(P_m)}$ under updated $\gamma^{(q)}$; $V_m^{(q)} \leftarrow \sum_i \gamma_{im}^{(q)} \|s^{(i)} - \bar{s}_m^{(q)}\|_w^2$
- 22: **if** $\sum_i \gamma_{im}^{(q)} \|r_m^{(P_m, i)}\|_w^2 \geq \tau \cdot V_m^{(q)}$ **then**
- 23: Train mode $P_m + 1$; $P_m \leftarrow P_m + 1$
- 24: **end if**
- 25: **if** $\|\phi_{P_m}^{(m)}\|_w^2 \cdot \sum_{i=1}^N \gamma_{im}^{(q)} (\lambda_{P_m}^{(m, i)})^2 < \varepsilon_p$ **then**
- 26: $P_m \leftarrow P_m - 1$
- 27: **end if**
- 28: **else** FULL RERUN
- 29: $\Phi_m^{(q)}, \Lambda_m^{(q)} \leftarrow \text{WEIGHTEDNEURALPOD}(\{s^{(i)}, \gamma_{im}^{(q)}\}, \tau)$
- 30: **end if**
- 31: **end for**
- 32: **until** $\max_m \Delta_m^{(q)} < \varepsilon_c$
- 33: **return** $\Phi_m^* \leftarrow \Phi_m^{(q)}, \Lambda_m^* \leftarrow \Lambda_m^{(q)}, \gamma_{im}^* \leftarrow \gamma_{im}^{(q)}, \pi_m^* \leftarrow \pi_m^{(q)}$
